



US008458163B2

(12) **United States Patent**  
**Fox et al.**

(10) **Patent No.:** **US 8,458,163 B2**  
(45) **Date of Patent:** **\*Jun. 4, 2013**

(54) **SYSTEM AND METHOD FOR ENABLING WEBSITE OWNER TO MANAGE CRAWL RATE IN A WEBSITE INDEXING SYSTEM**

USPC ..... 707/709; 707/E17.108; 707/802; 709/224; 709/226

(75) Inventors: **Vanessa Fox**, Redmond, WA (US); **Amanda Ann Camp**, Kirkland, WA (US); **Maximilian Ibel**, Pfaffikon (CH); **Patrik Rene Celeste Reali**, Zurich (CH); **Jeremy J. Lilley**, Mountain View, CA (US); **Katherine Jane Lai**, Cambridge, MA (US); **Ted J. Bonkenburg**, Mountain View, CA (US); **Neal Douglas Cardwell**, San Francisco, CA (US)

(58) **Field of Classification Search**

CPC ..... G06F 17/30887  
USPC ..... 707/709, 713; 709/224, 225, 226  
See application file for complete search history.

(73) Assignee: **Google Inc.**, Mountain View, CA (US)

(\*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 0 days.

This patent is subject to a terminal disclaimer.

(56) **References Cited**

U.S. PATENT DOCUMENTS

5,935,210	A	8/1999	Stark
5,958,008	A	9/1999	Pogrebisky et al.
6,124,966	A	9/2000	Yokoyama
6,144,959	A	11/2000	Anderson et al.
6,269,370	B1	7/2001	Kirsch
6,271,840	B1	8/2001	Finseth et al.
6,285,999	B1	9/2001	Page

(Continued)

OTHER PUBLICATIONS

Dikaiakos et al., K. Bauknecht, A Min Tjoa, G. Quirchmayr (Eds.): EC-Web 2003, LNCS 2738, pp. 369-378, 2003.\*

(21) Appl. No.: **13/252,149**

(Continued)

(22) Filed: **Oct. 3, 2011**

(65) **Prior Publication Data**

US 2012/0023091 A1 Jan. 26, 2012

**Related U.S. Application Data**

(63) Continuation of application No. 12/554,757, filed on Sep. 4, 2009, now Pat. No. 8,032,518, which is a continuation of application No. 11/549,075, filed on Oct. 12, 2006, now Pat. No. 7,599,920.

(51) **Int. Cl.**  
**G06F 17/30** (2006.01)

(52) **U.S. Cl.**  
CPC .... **G06F 17/30864** (2013.01); **G06F 17/30887** (2013.01); **G06F 17/3089** (2013.01)

*Primary Examiner* — Cheyne Ly

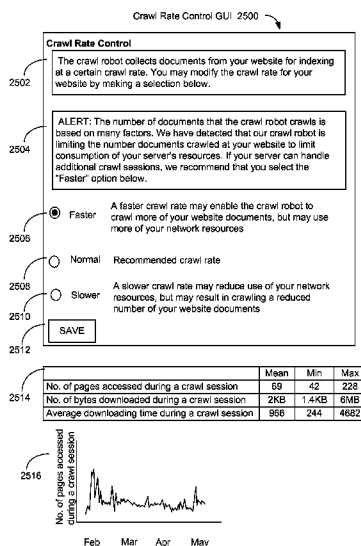
(74) *Attorney, Agent, or Firm* — Morgan, Lewis & Bockius LLP

(57)

**ABSTRACT**

Web crawlers crawl websites to access documents of the website for purposes of indexing the documents for search engines. The web crawlers crawl a specified website at a crawl rate that is based on multiple factors. One of the factors is a pre-set crawl rate limit. According to certain embodiments, an owner for a specified website is enabled to modify the crawl rate limit for the specified website when one or more pre-set criteria are met.

**39 Claims, 23 Drawing Sheets**



## U.S. PATENT DOCUMENTS

6,321,265 B1 11/2001 Najork et al.  
 6,418,433 B1 7/2002 Chakrabarti et al.  
 6,421,724 B1 7/2002 Nickerson et al.  
 6,424,966 B1 7/2002 Meyerzon et al.  
 6,516,337 B1 2/2003 Tripp et al.  
 6,525,748 B1 2/2003 Belfiore et al.  
 6,636,854 B2 10/2003 Dutta et al.  
 6,732,105 B1 5/2004 Watson, Jr. et al.  
 6,957,383 B1 10/2005 Smith  
 6,976,053 B1 12/2005 Tripp et al.  
 6,983,282 B2 1/2006 Stern et al.  
 7,133,870 B1 11/2006 Tripp et al.  
 7,139,747 B1 11/2006 Najork  
 7,191,210 B2 3/2007 Grossman  
 7,769,742 B1 \* 8/2010 Brawer et al. .... 707/709  
 7,774,782 B1 8/2010 Popescu et al.  
 7,844,610 B2 11/2010 Hillis et al.  
 7,987,172 B1 \* 7/2011 Carver ..... 707/709  
 2002/0032772 A1 \* 3/2002 Olstad et al. .... 709/224  
 2002/0052928 A1 5/2002 Stern et al.  
 2002/0061029 A1 5/2002 Dillon  
 2002/0087515 A1 7/2002 Swannack et al.  
 2002/0138582 A1 9/2002 Chandra et al.  
 2003/0028896 A1 2/2003 Swart et al.  
 2003/0208482 A1 \* 11/2003 Kim et al. .... 707/3  
 2004/0030683 A1 2/2004 Evans et al.  
 2004/0093327 A1 5/2004 Anderson et al.  
 2004/0122686 A1 6/2004 Hill et al.  
 2004/0143787 A1 \* 7/2004 Grancharov et al. .... 715/501.1  
 2004/0158617 A1 8/2004 Shanny et al.  
 2004/0168066 A1 8/2004 Alden  
 2004/0221289 A1 11/2004 D'Souza et al.  
 2005/0060286 A1 \* 3/2005 Hansen et al. .... 707/2  
 2005/0138435 A1 6/2005 Kaufman et al.  
 2005/0256865 A1 11/2005 Ma et al.  
 2006/0004691 A1 1/2006 Sifry  
 2006/0070022 A1 3/2006 Ng et al.

2006/0080405 A1 4/2006 Gibson  
 2006/0106866 A1 5/2006 Green et al.  
 2006/0212451 A1 9/2006 Serdy, Jr. et al.  
 2007/0011168 A1 1/2007 Keohane et al.  
 2007/0050338 A1 \* 3/2007 Strohm et al. .... 707/3  
 2007/0067395 A1 3/2007 Blinn et al.  
 2008/0021904 A1 1/2008 Garg et al.

## OTHER PUBLICATIONS

Clabo, S., Hermetic Sitemap Builder, A Search Engine Submission and Website Development Tool, <http://www.hermetic.ch/smb/smb.htm>, Jan. 7, 2010, 3 pages.  
 Sitemaps.org, "What are Sitemaps?," <http://web.archive.org/web/20061205095241/http://sitemaps.org/>, Nov. 16, 2006, 1 page.  
 XML-Sitemaps.com, "Build Your Site Map Online," <http://www.xml-sitemaps.com>, © 2005, 2 pages.  
 Yahoo! Adopts Site Maps (urllist.txt), "The Optimizer—Weekly SEO News," <http://www.increased-online-traffic.com/2005/08/yahoo-adopts-site-maps-urllisttxt.asp>, Aug. 23, 2005, 5 pages.  
 Yahoo! Free Sitemaps: Submit Your Site, <http://www.seroundtable.com/archives/002421.html>, Aug. 23, 2005, 2 pages.  
 Yahoo Search Blog, "Archive for the 'Site Explorer' Category," <http://www.ysearchblog.com/category/site-explorer/page/2/>, Dec. 6, 2005, Jun. 2, 2006 and Feb. 26, 2007, 11 pages.  
 Yahoo! Sitemap Feed Submission, Antezeta, Seo & Wab Analytics Consulting, <http://www.antezeta.com/yahoo/site-map-feed.html>, Feb. 26, 2010, 7 pages.  
 "Technorati: Ping Configurations", <http://web.archive.org/web/20040829035832/www.technorati.com/de...>, Aug. 2004.  
 "SOAP Meets RSS," <http://blogs.law.harvard.edu/tech/soapMeetsRss>, Jul. 17, 2003.  
 "The Open Archives Initiative Protocol for Metadata Harvesting," Ver. 2.0, <http://www.openarchives.org/OA/openarchivesprotocol.html>, Jun. 14, 2002.

\* cited by examiner

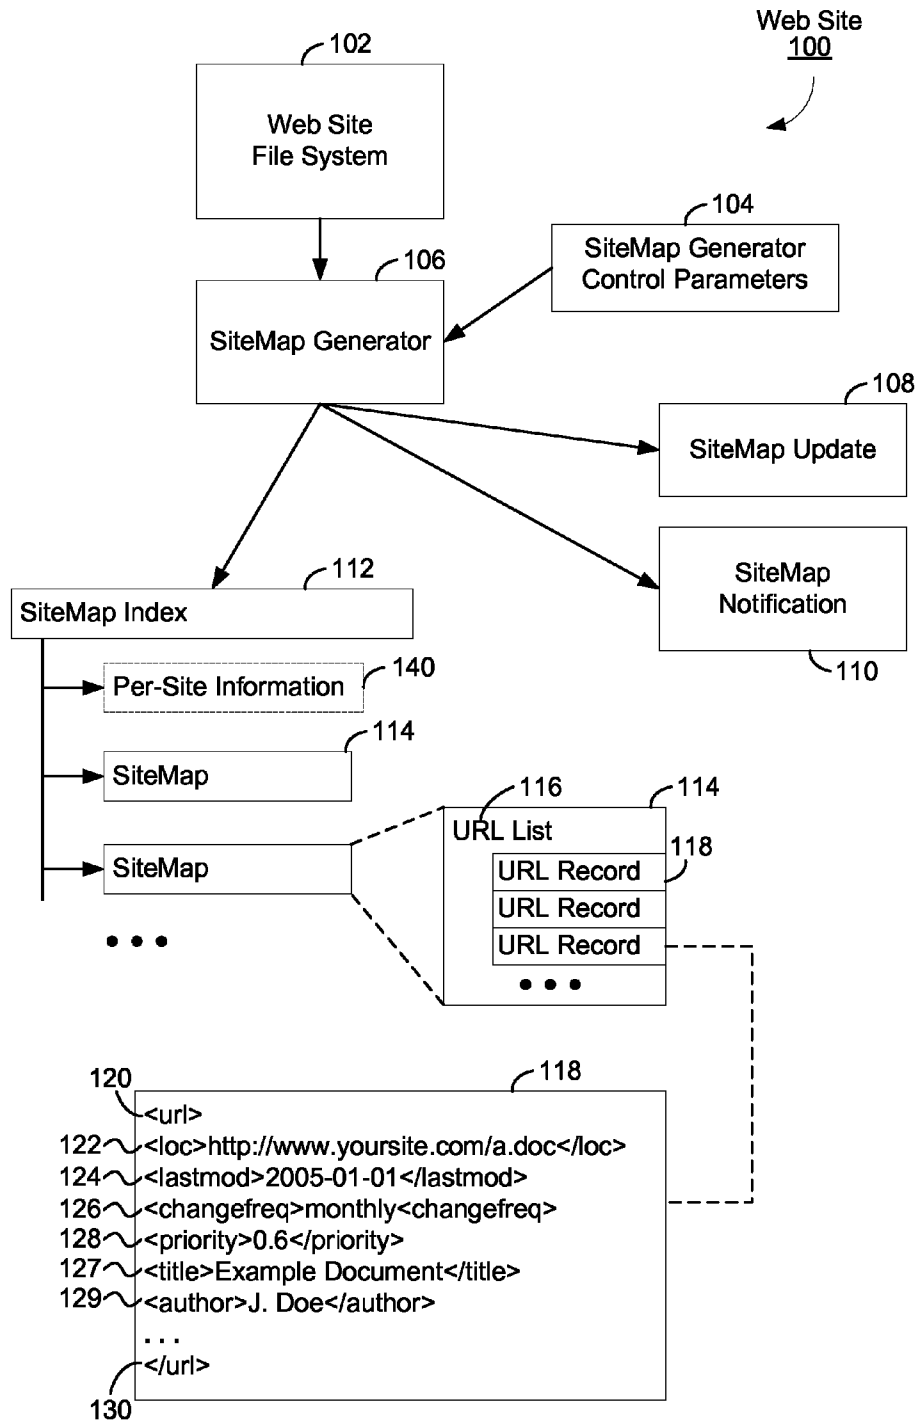


Figure 1

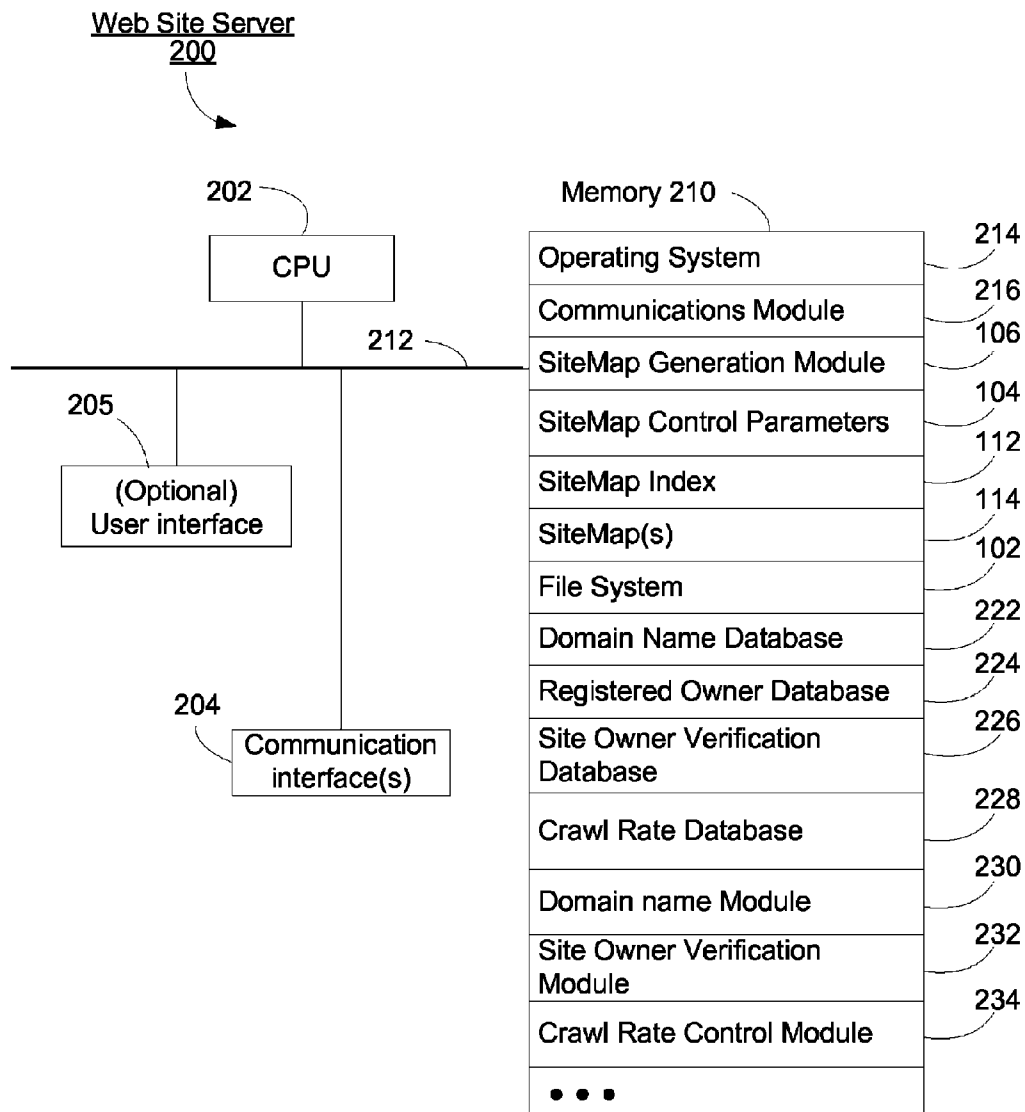


Figure 2

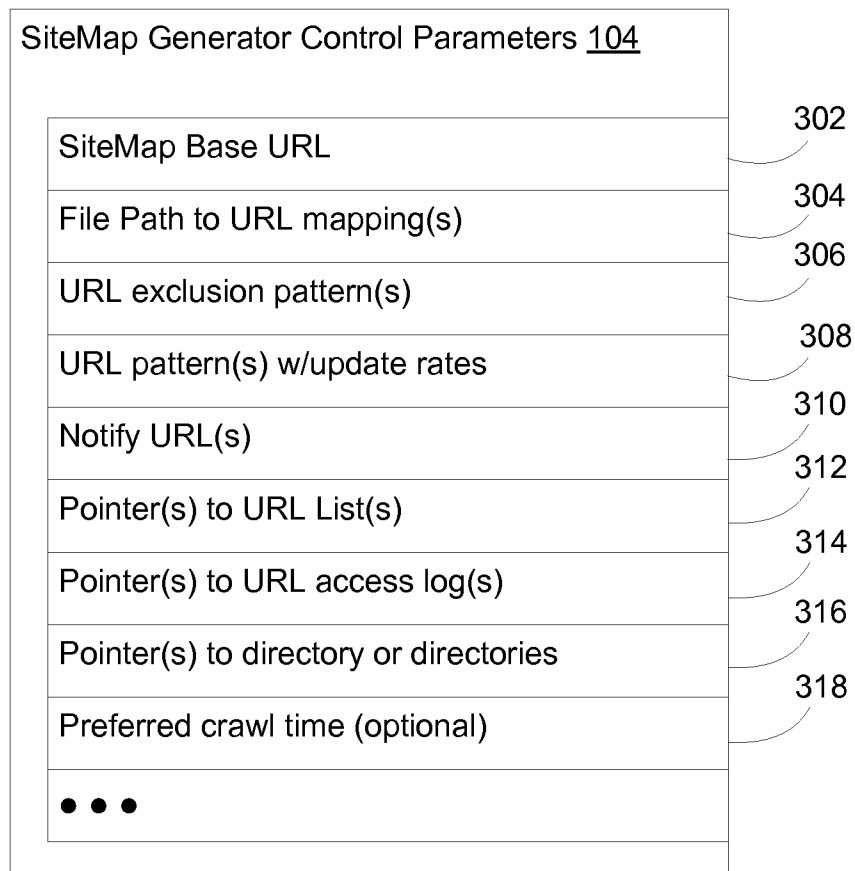


Figure 3

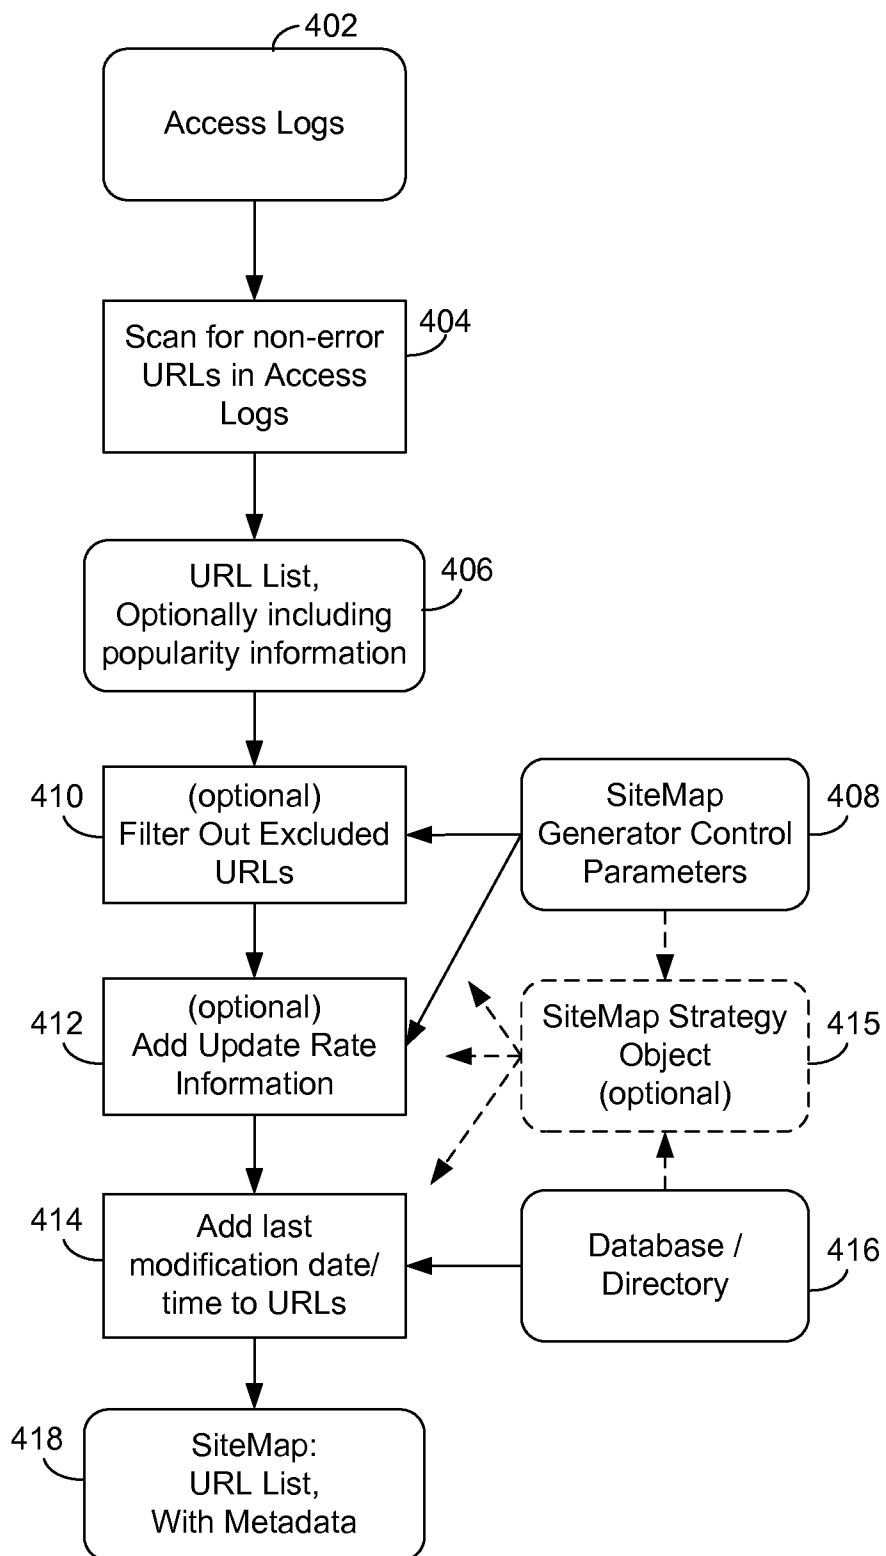


Figure 4

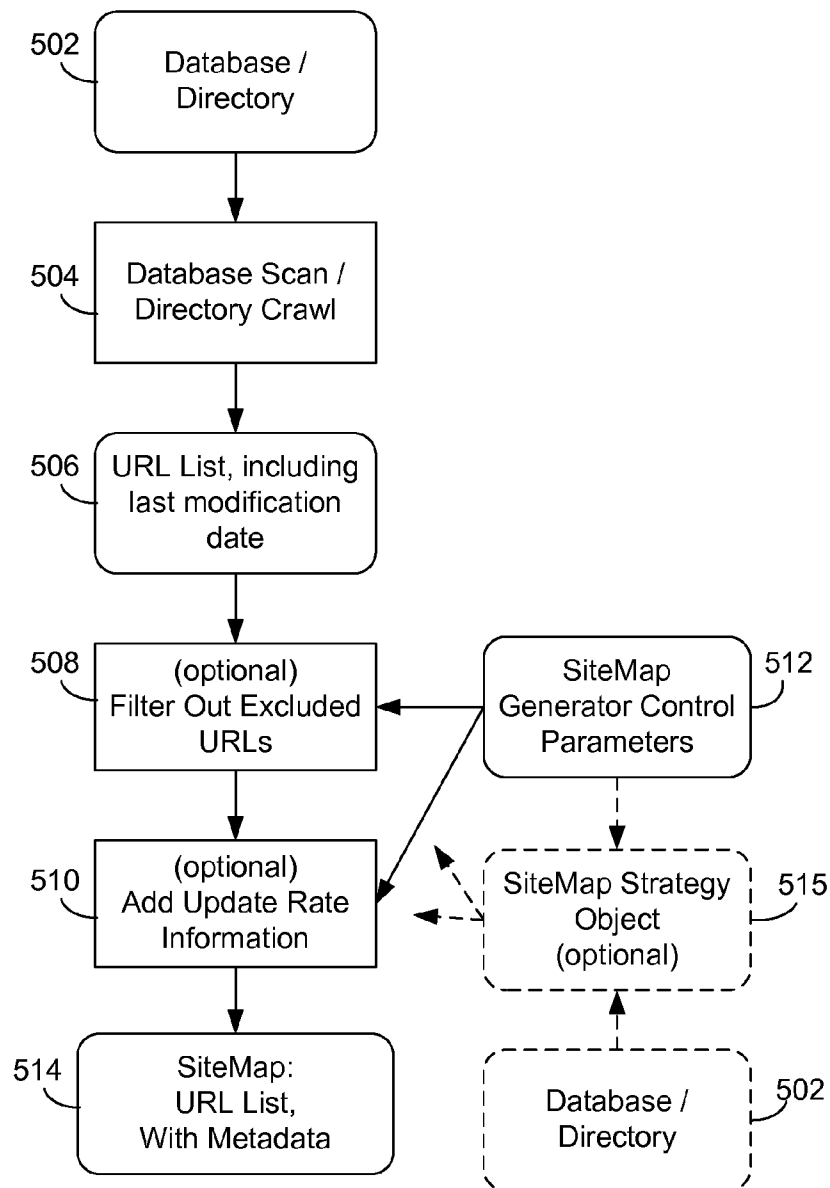


Figure 5

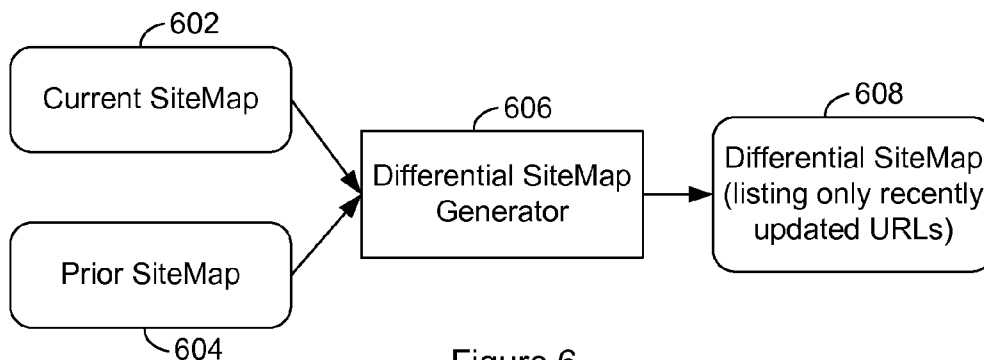


Figure 6

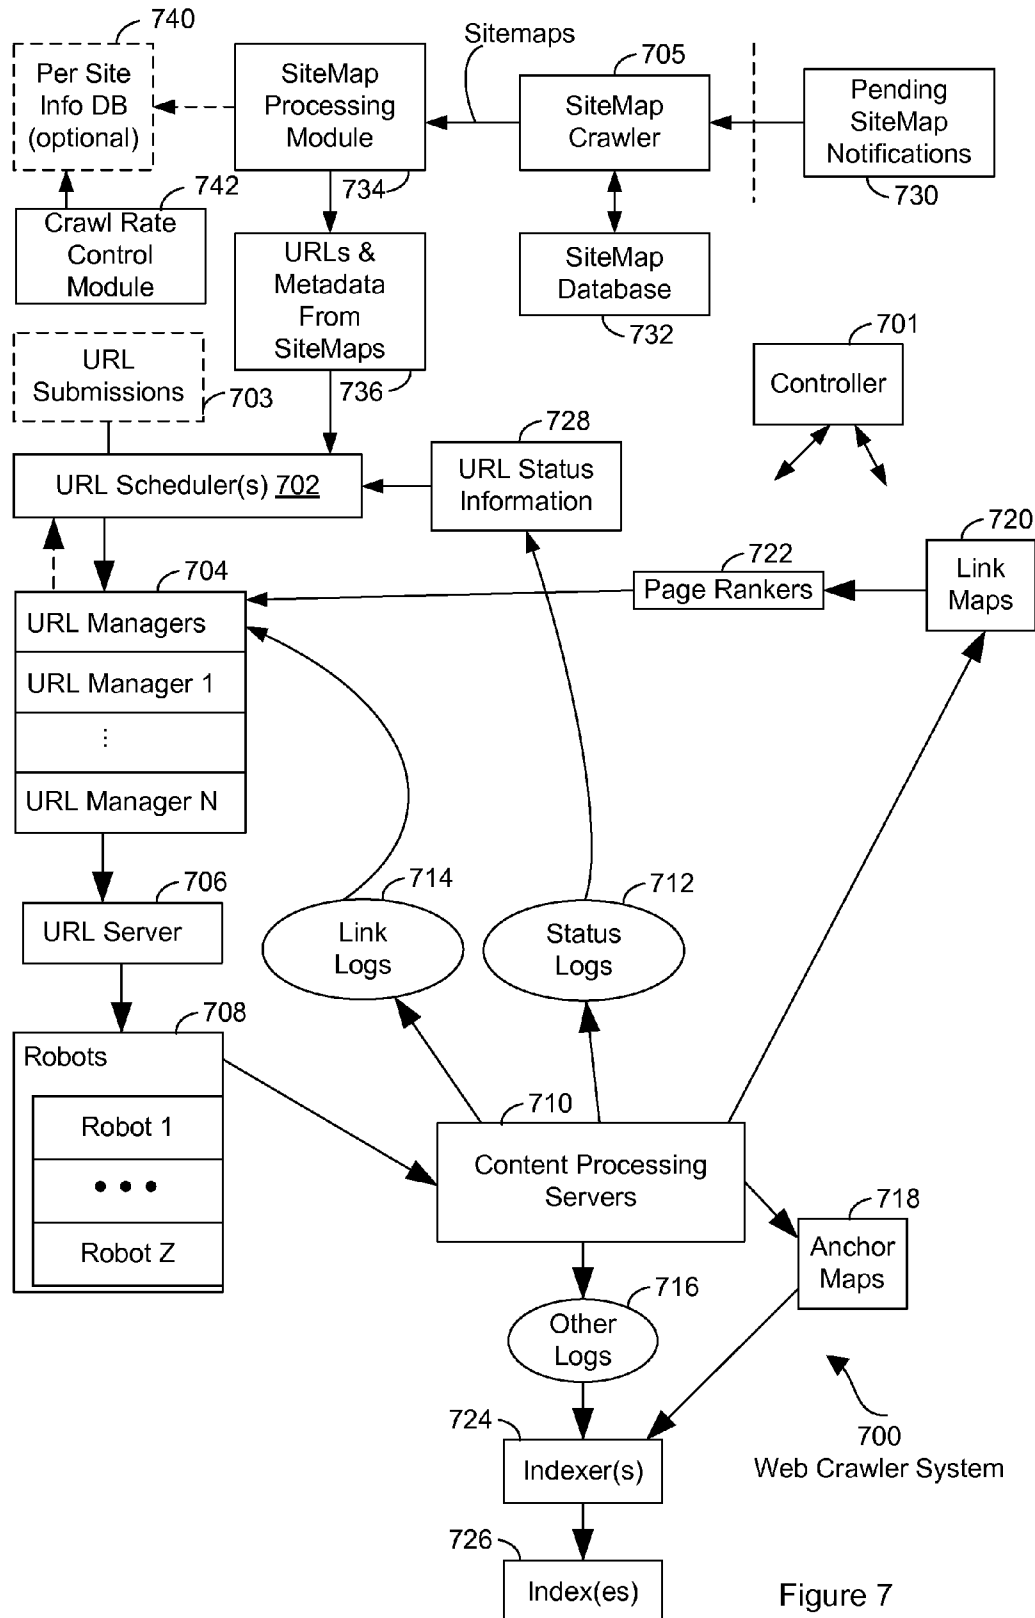


Figure 7



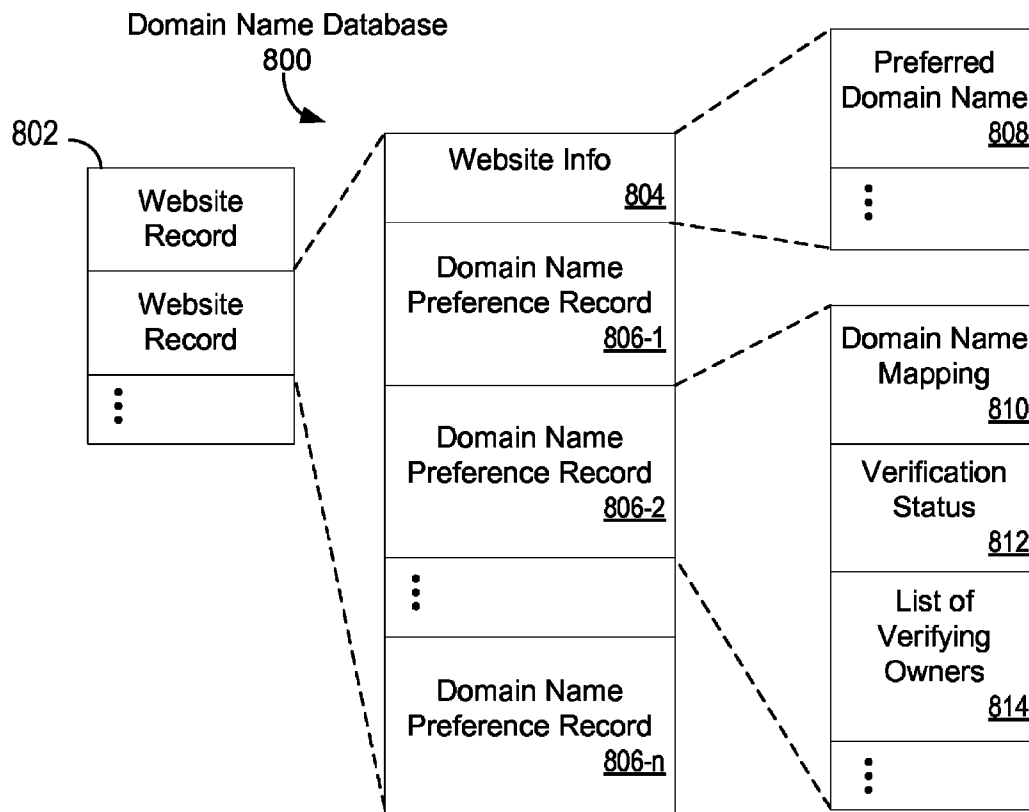


Figure 8

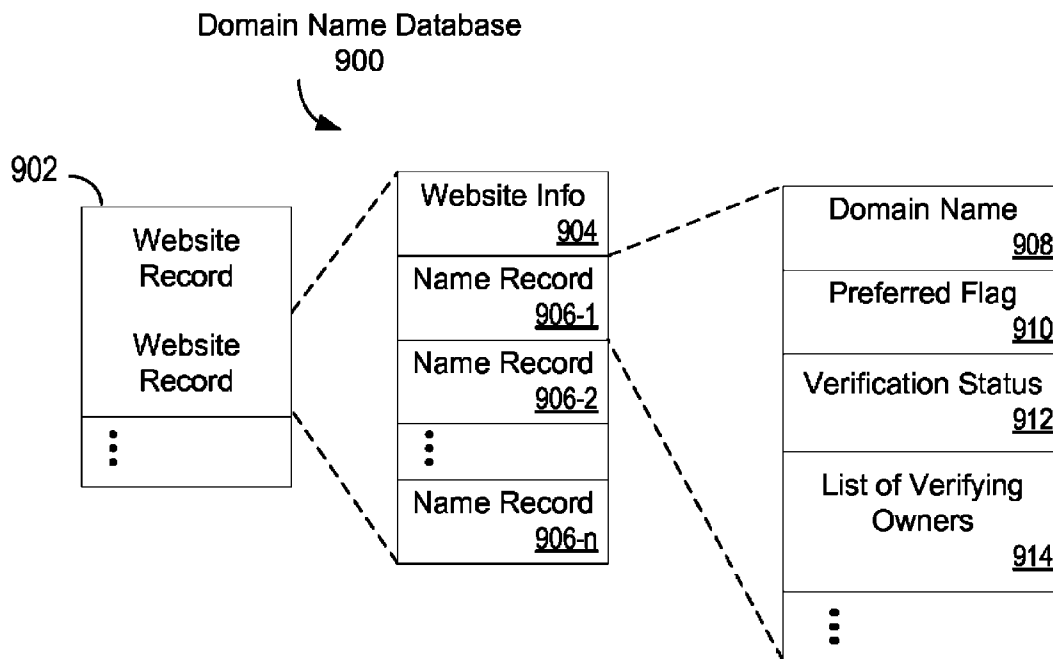


Figure 9

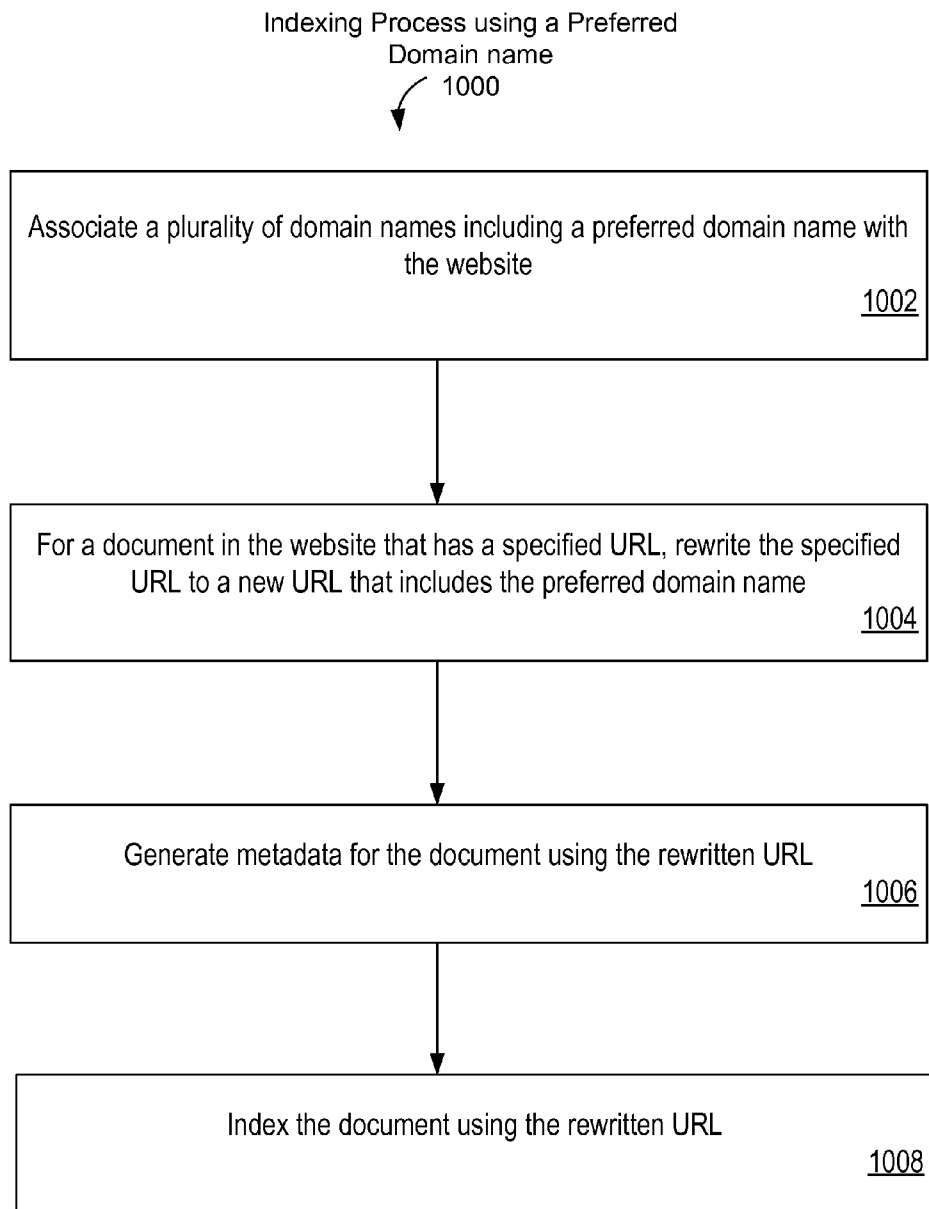


Figure 10

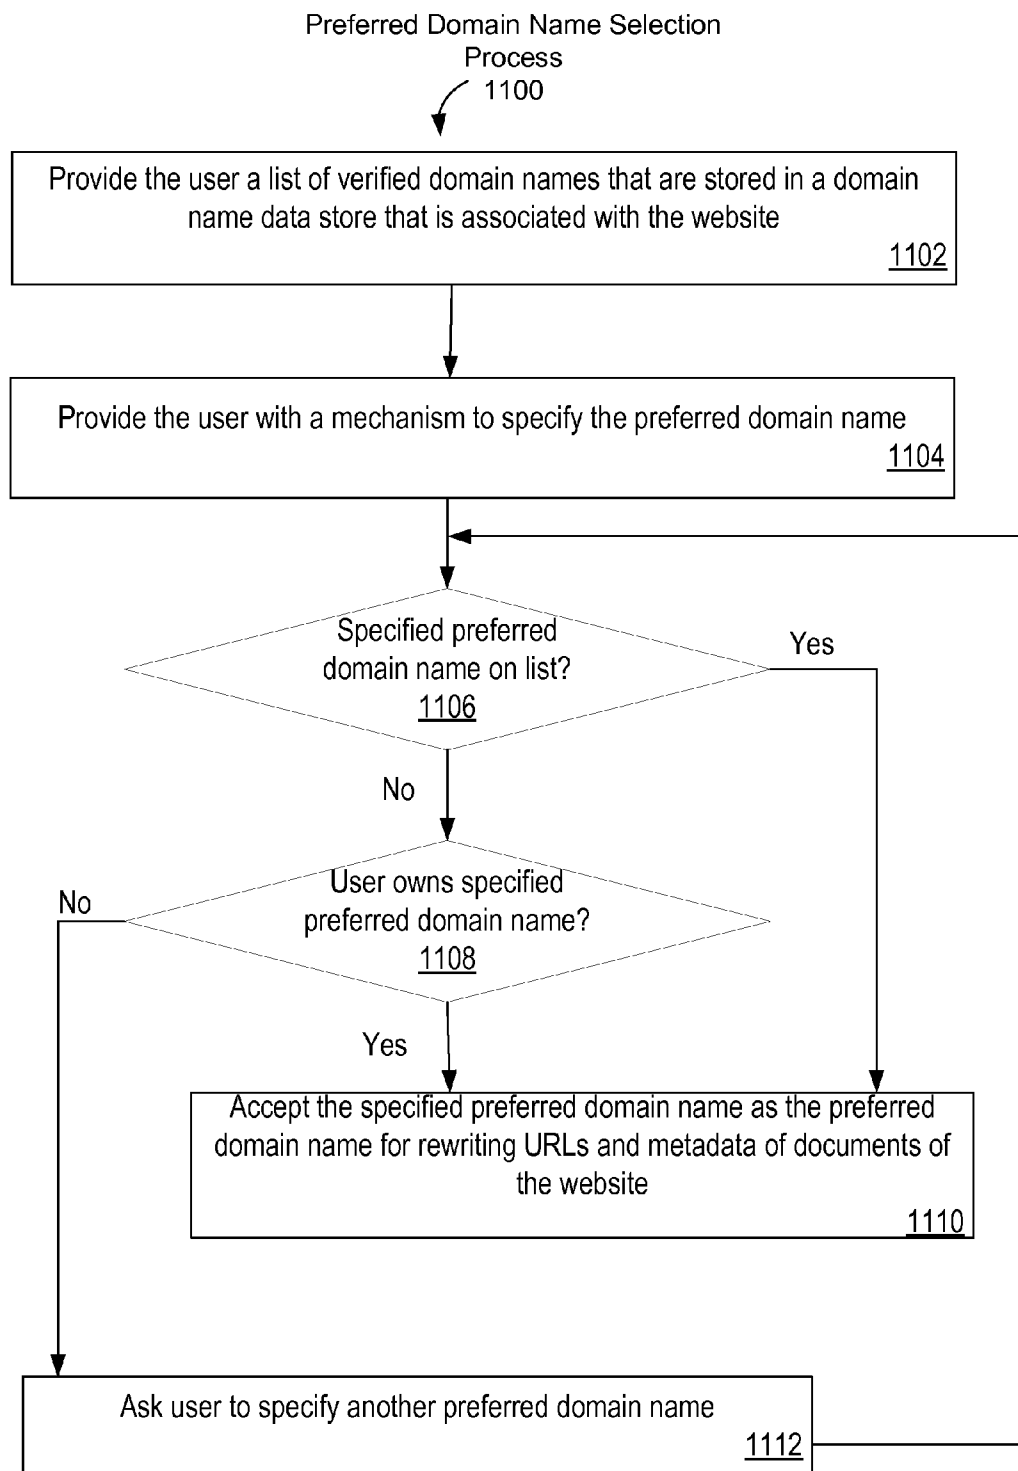


Figure 11

Preferred Domain GUI 1200

Preferred Domain Name [?]

1202 You may indicate the preferred domain name for indexing your website

1204 Display URLs as "www.domain.com" (for both "www.domain.com" and "domain.com")

1206 Display URLs as "domain.com" (for both "www.domain.com" and "domain.com")

1208 Do not set an association

1210 OK

Figure 12

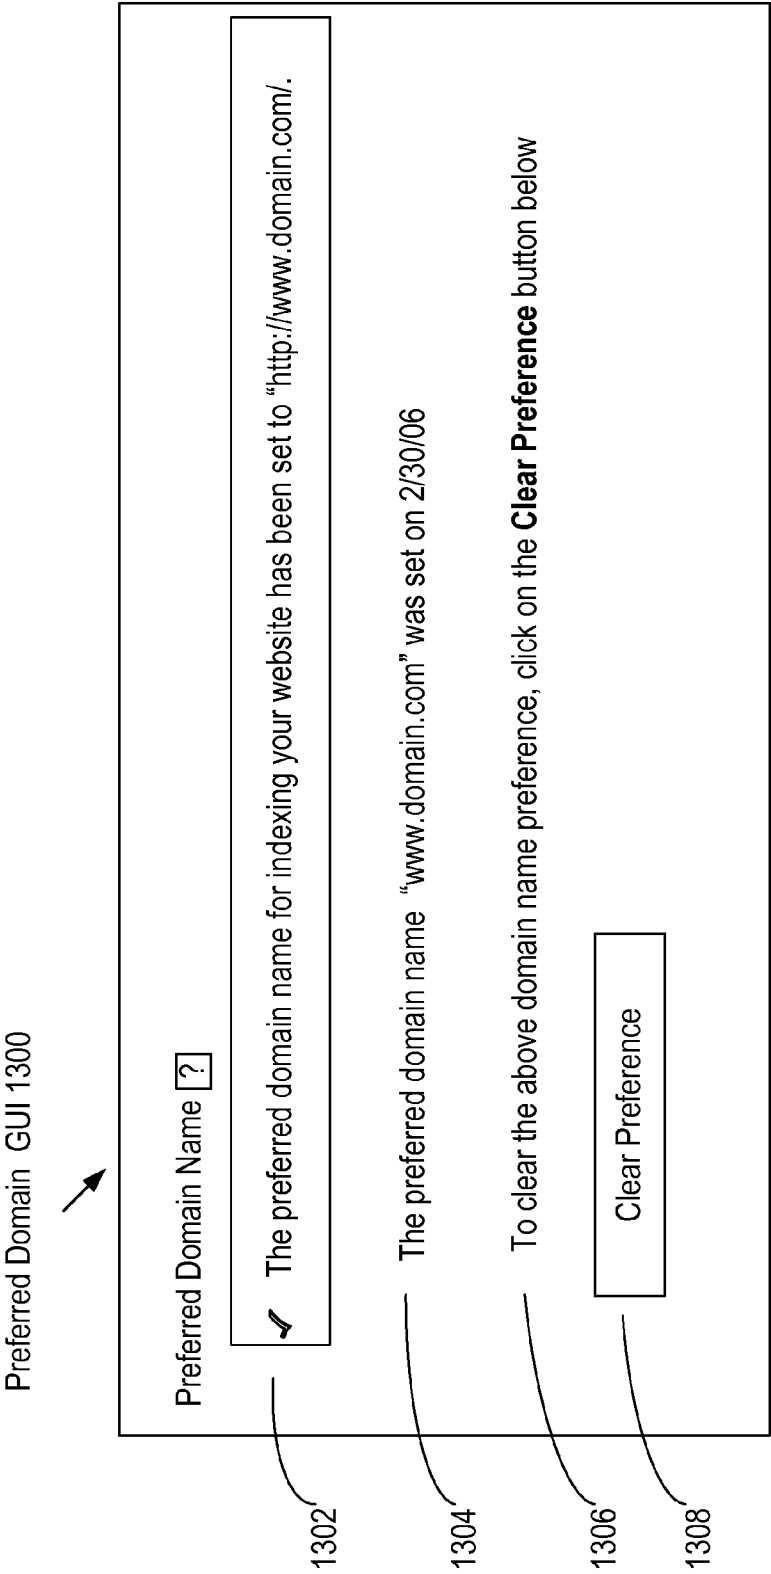
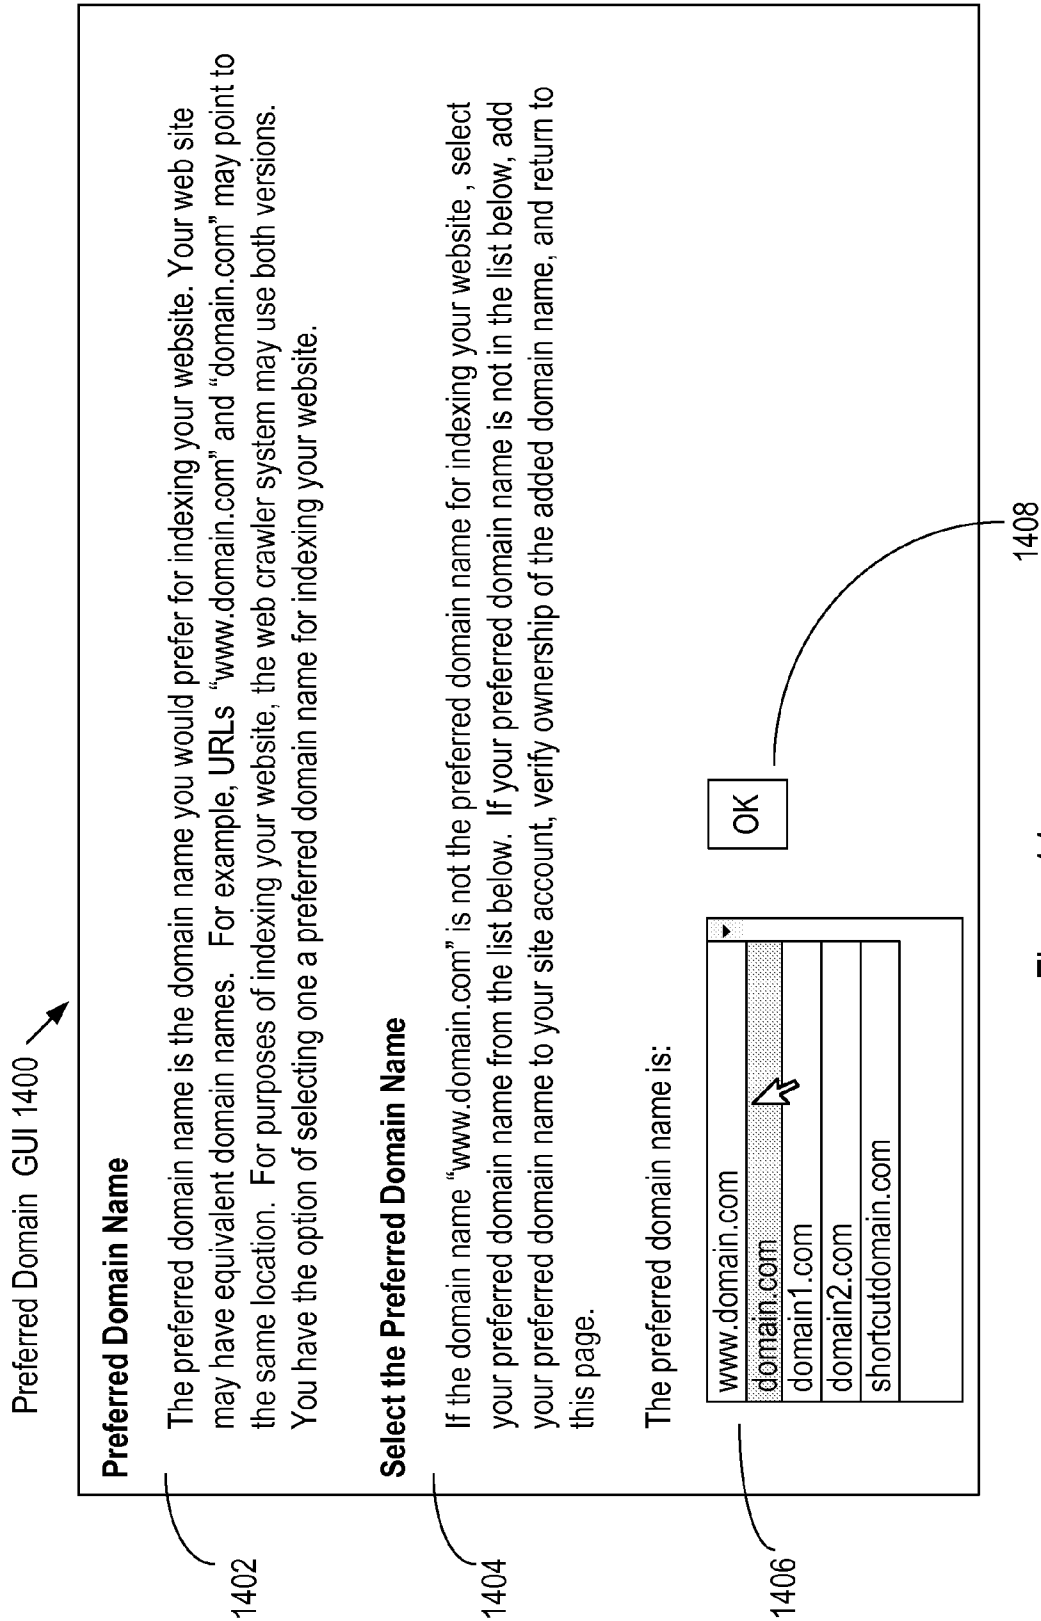


Figure 13



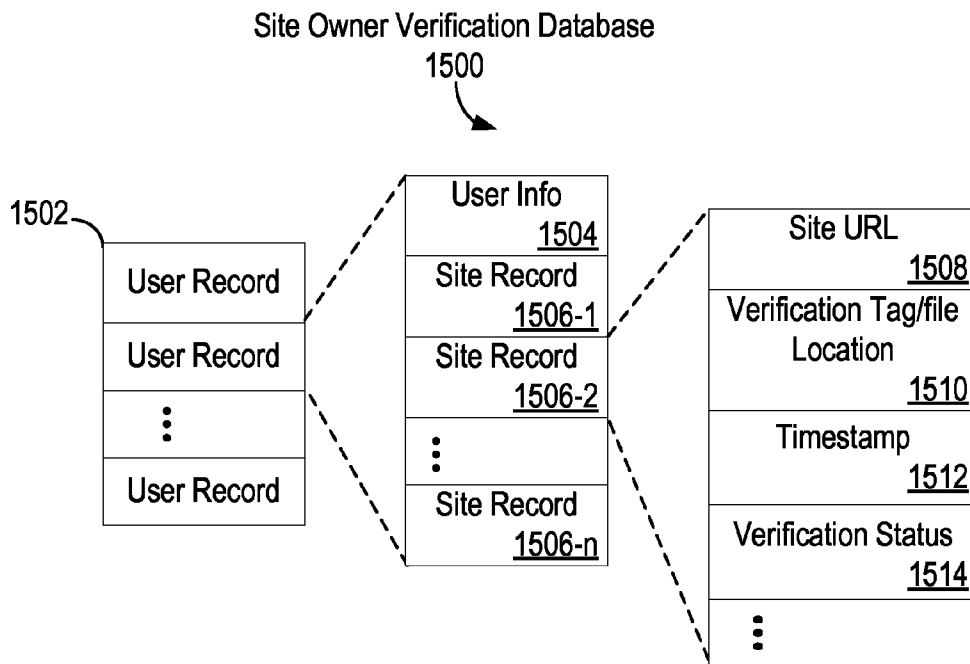


Figure 15

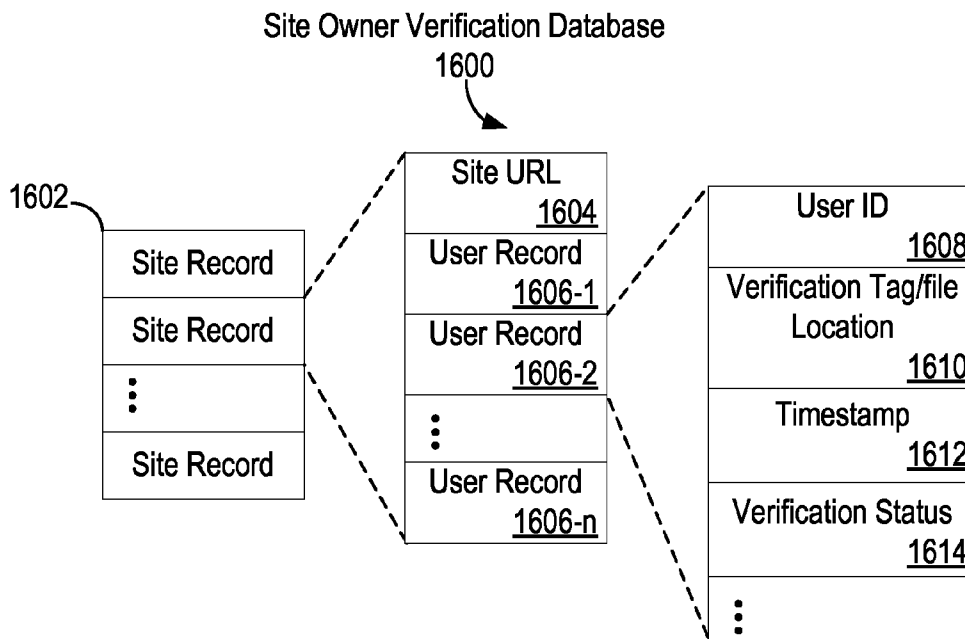


Figure 16



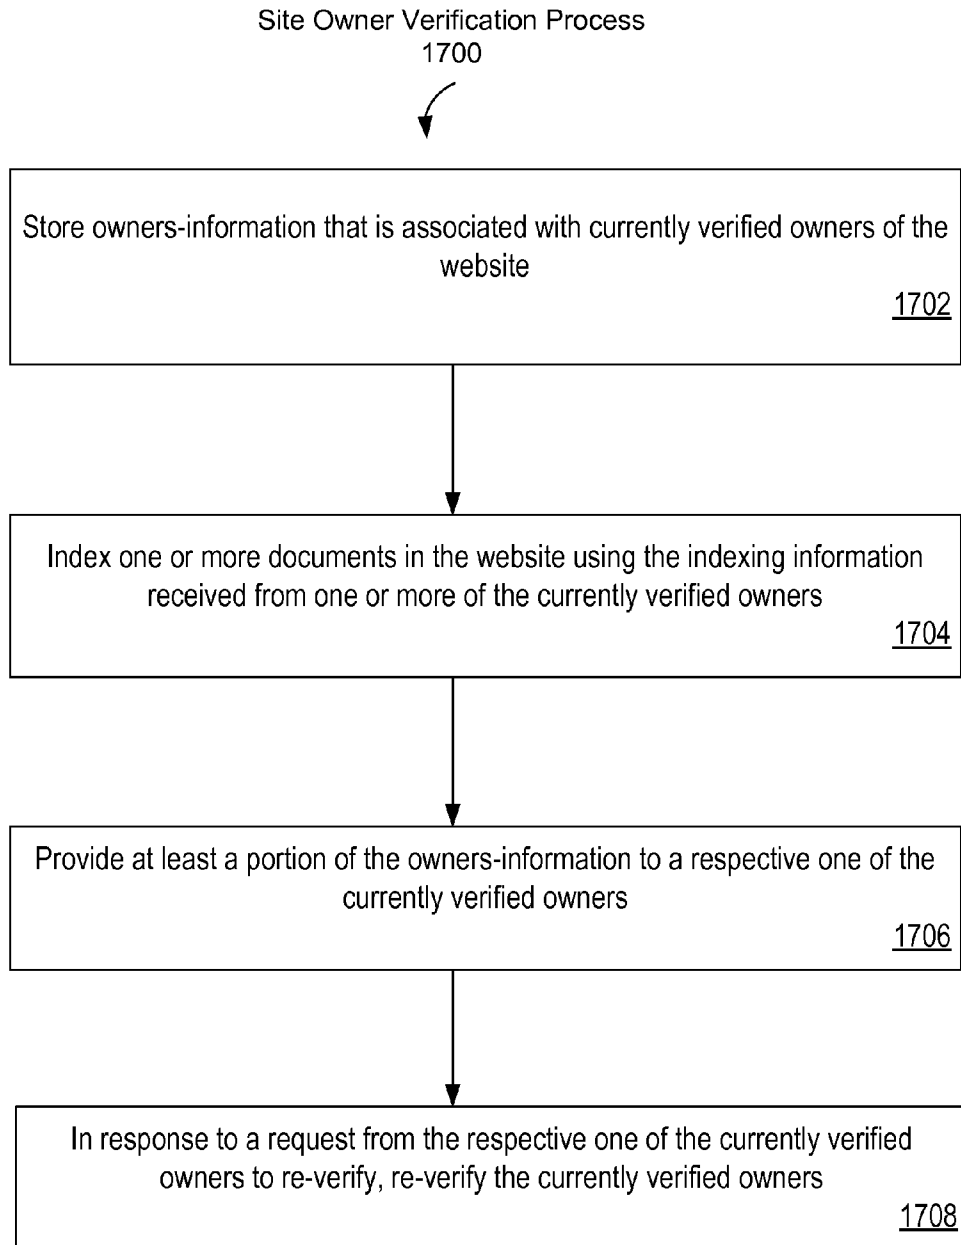


Figure 17

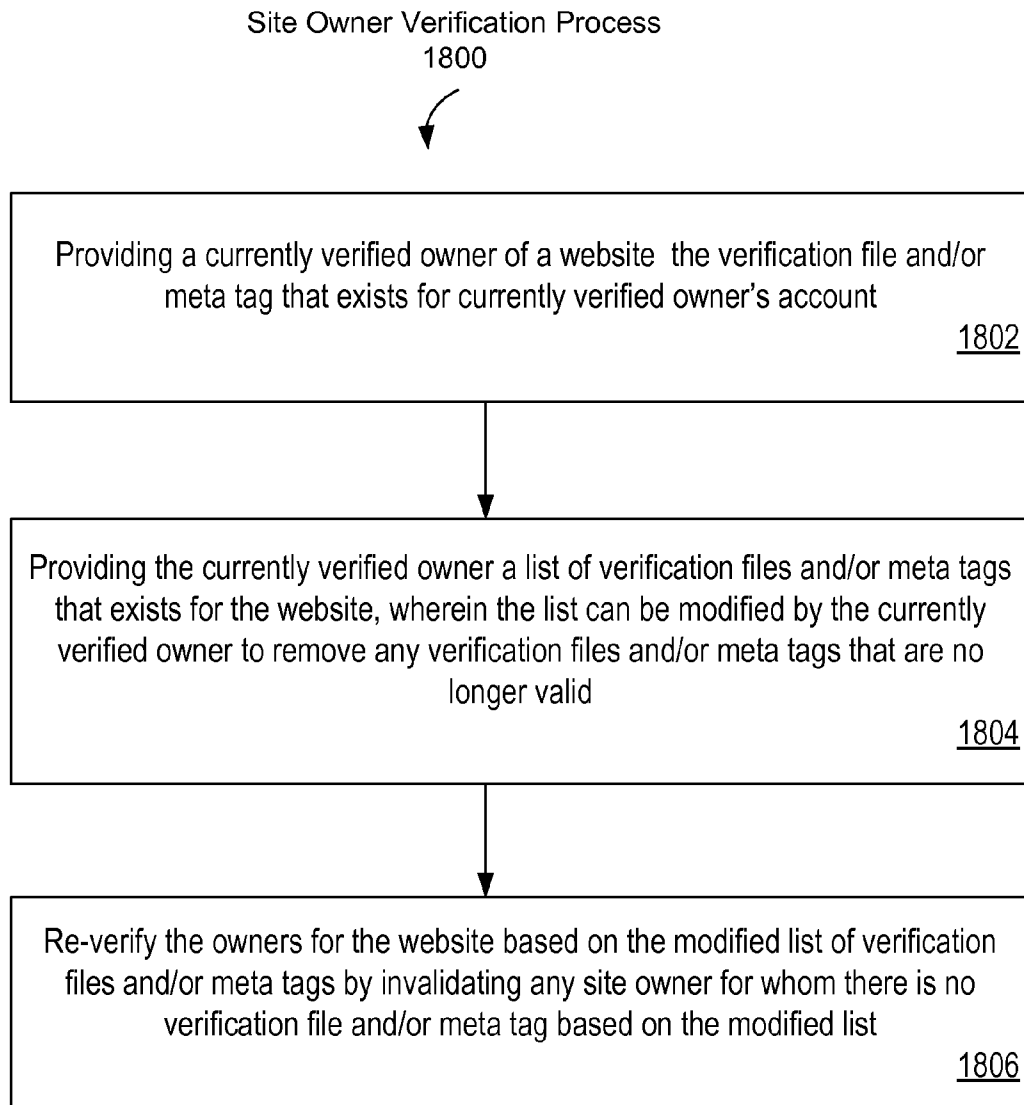


Figure 18

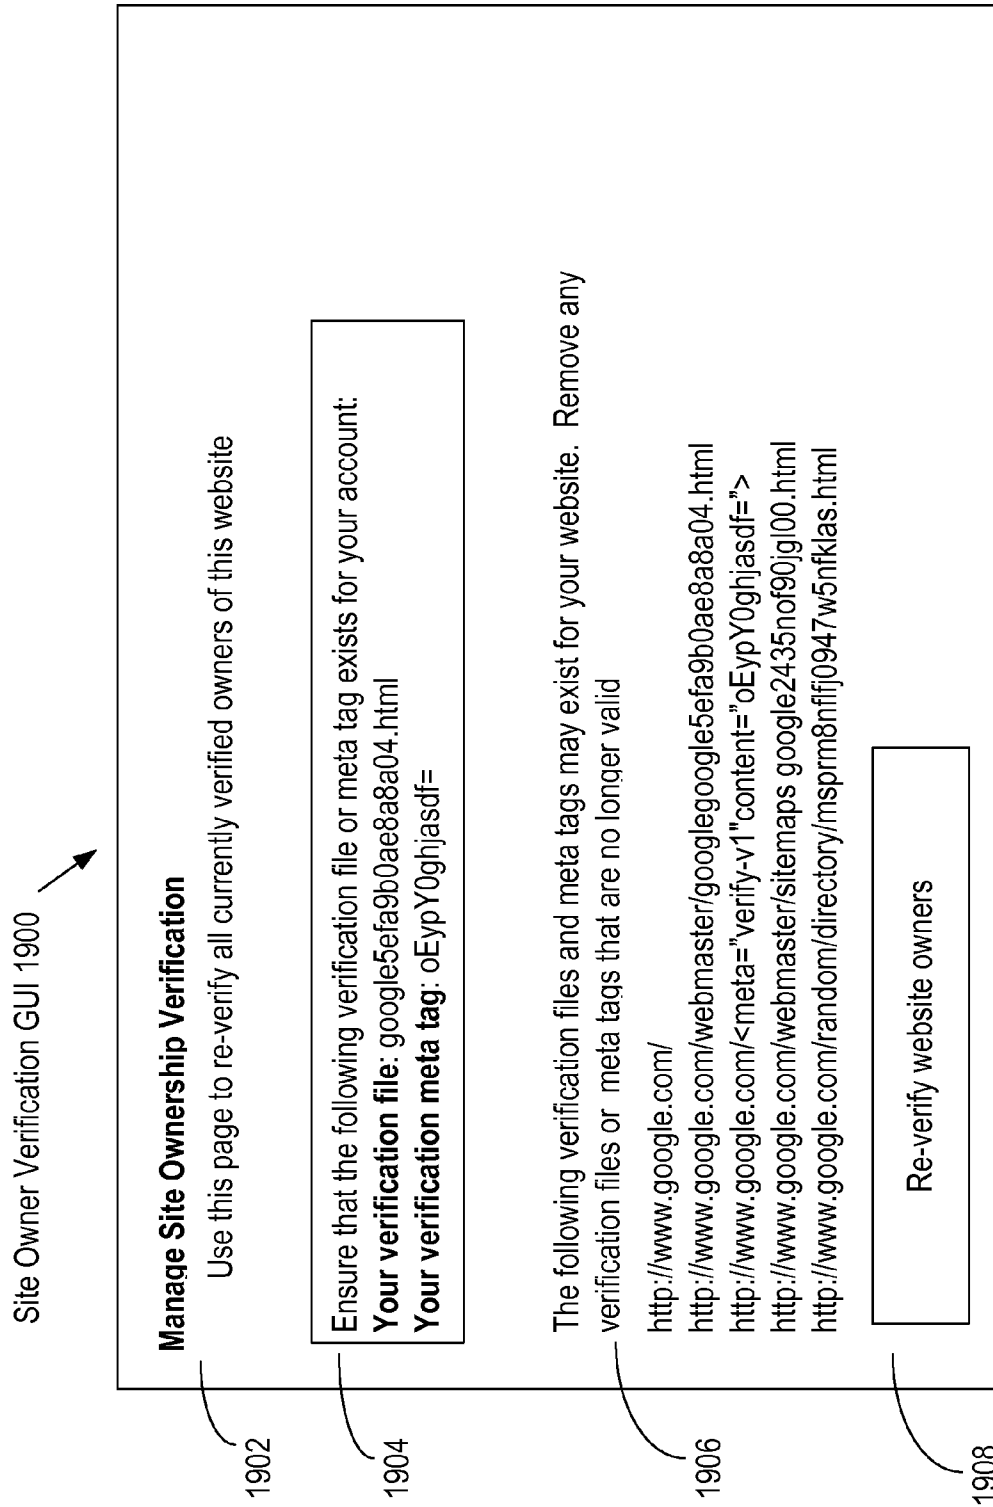


Figure 19

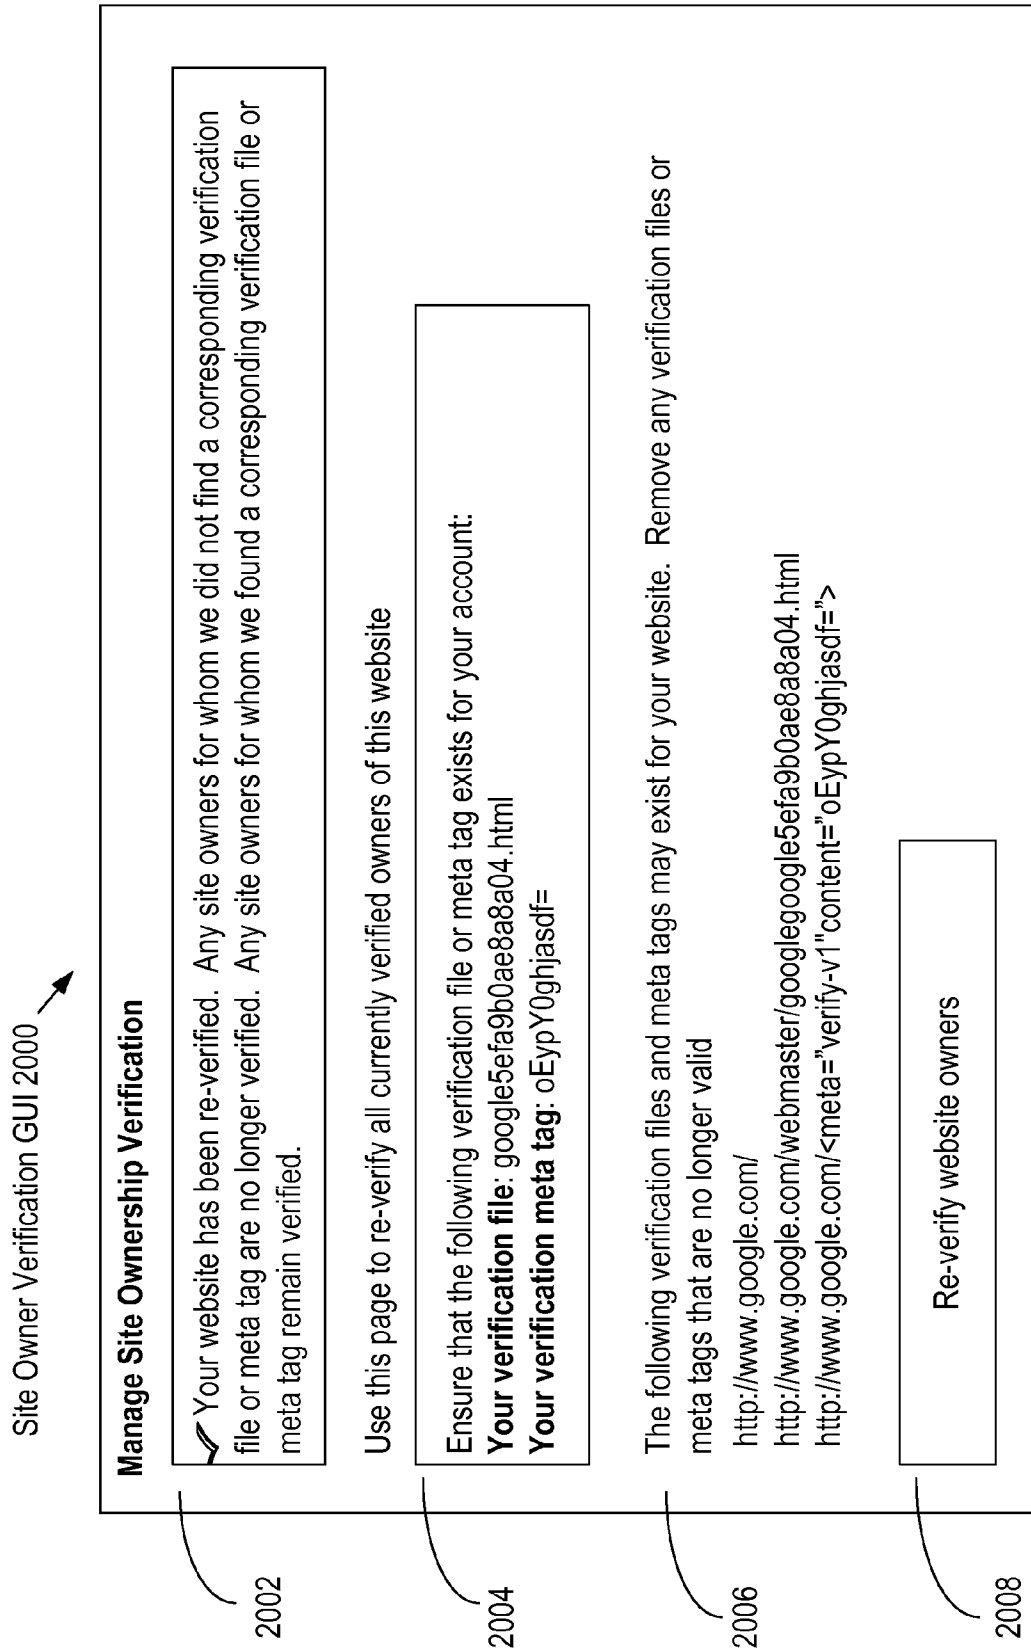


Figure 20

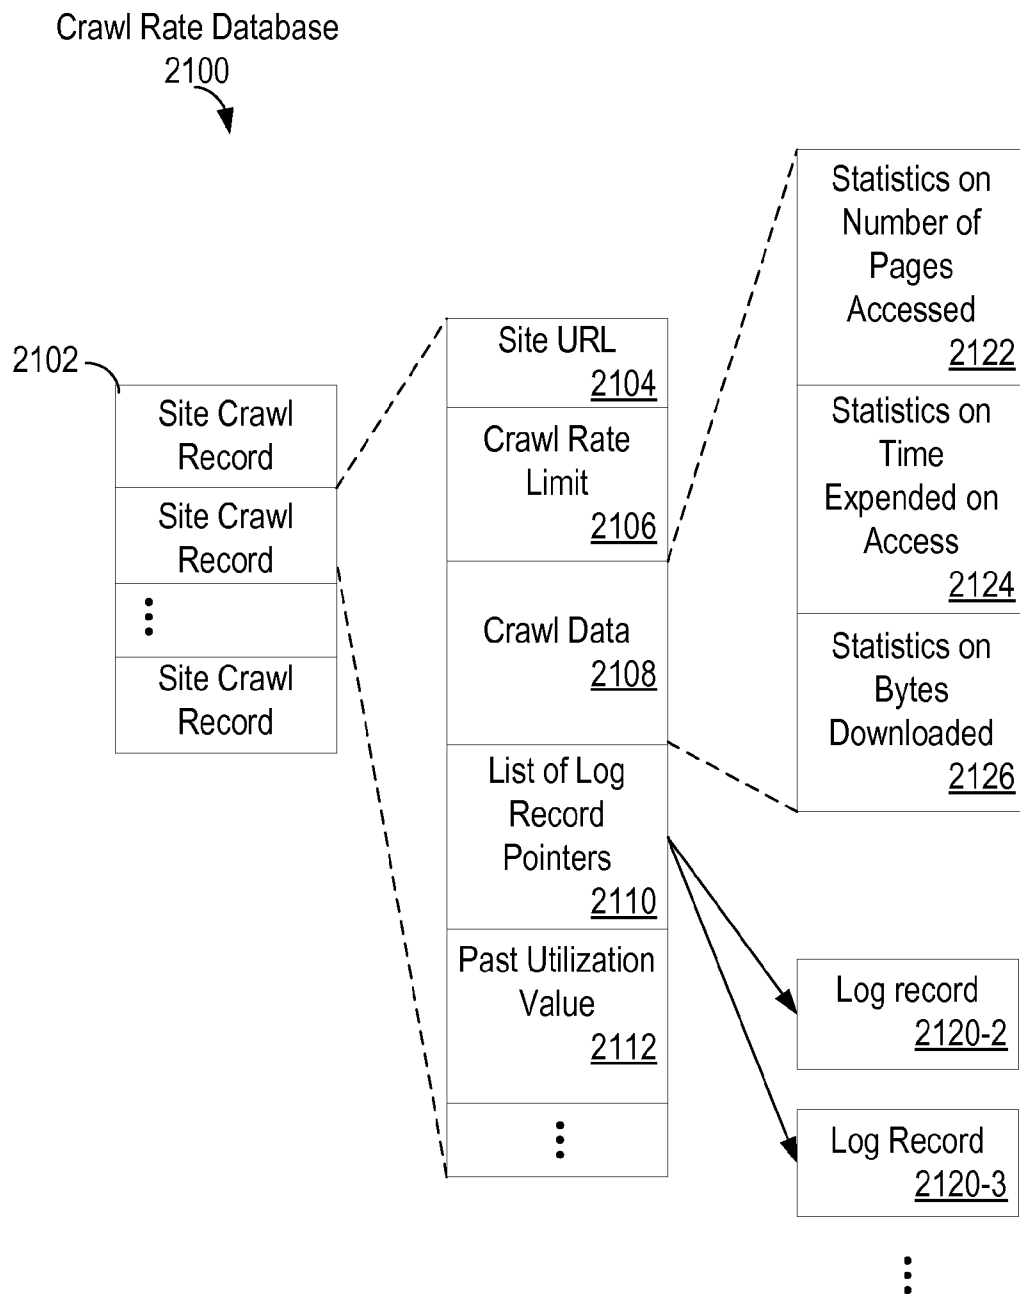


Figure 21

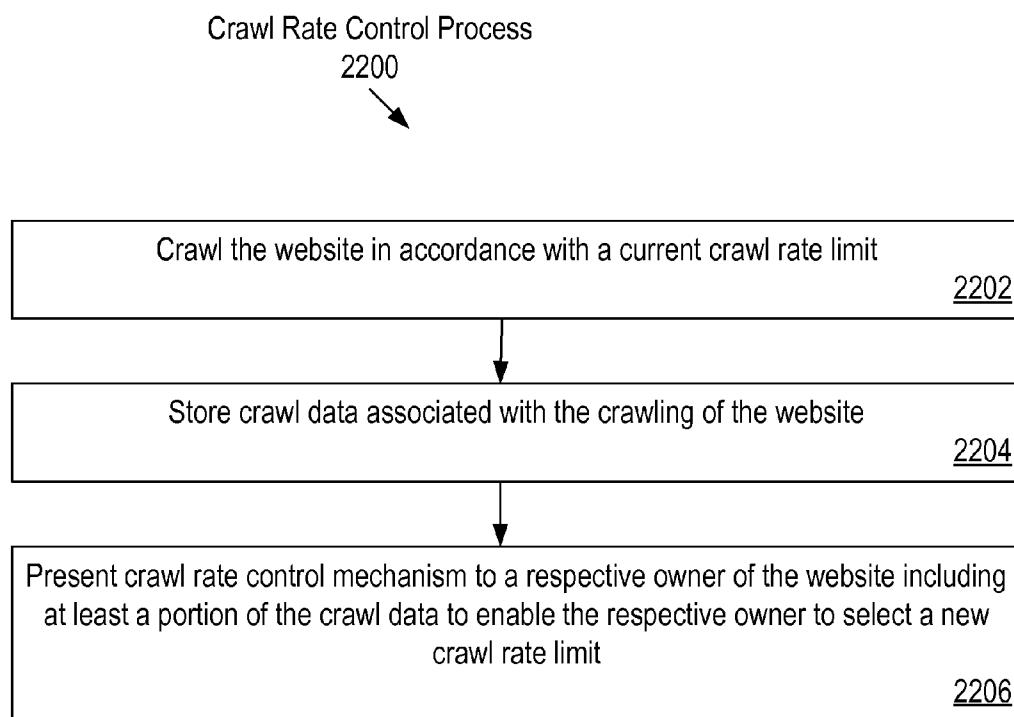


Figure 22

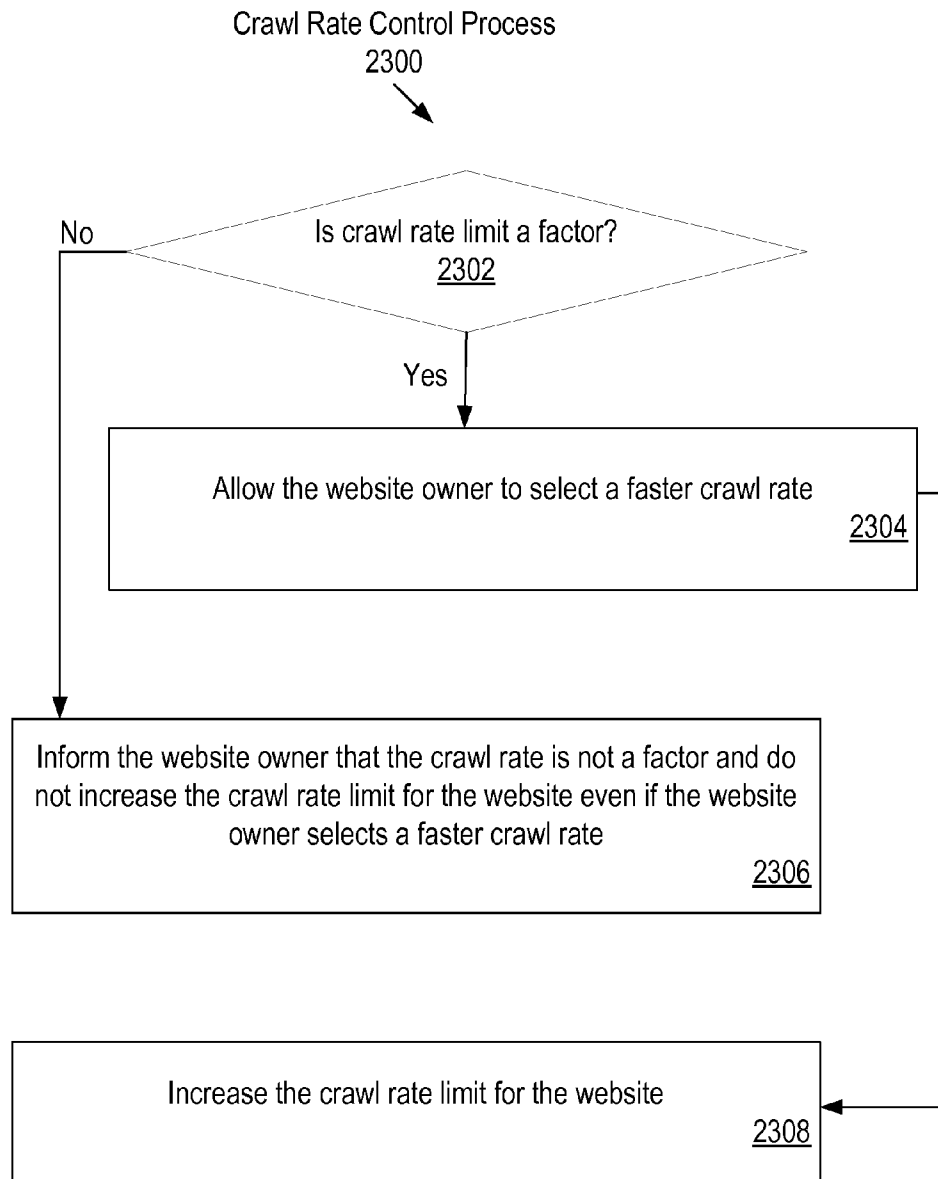


Figure 23

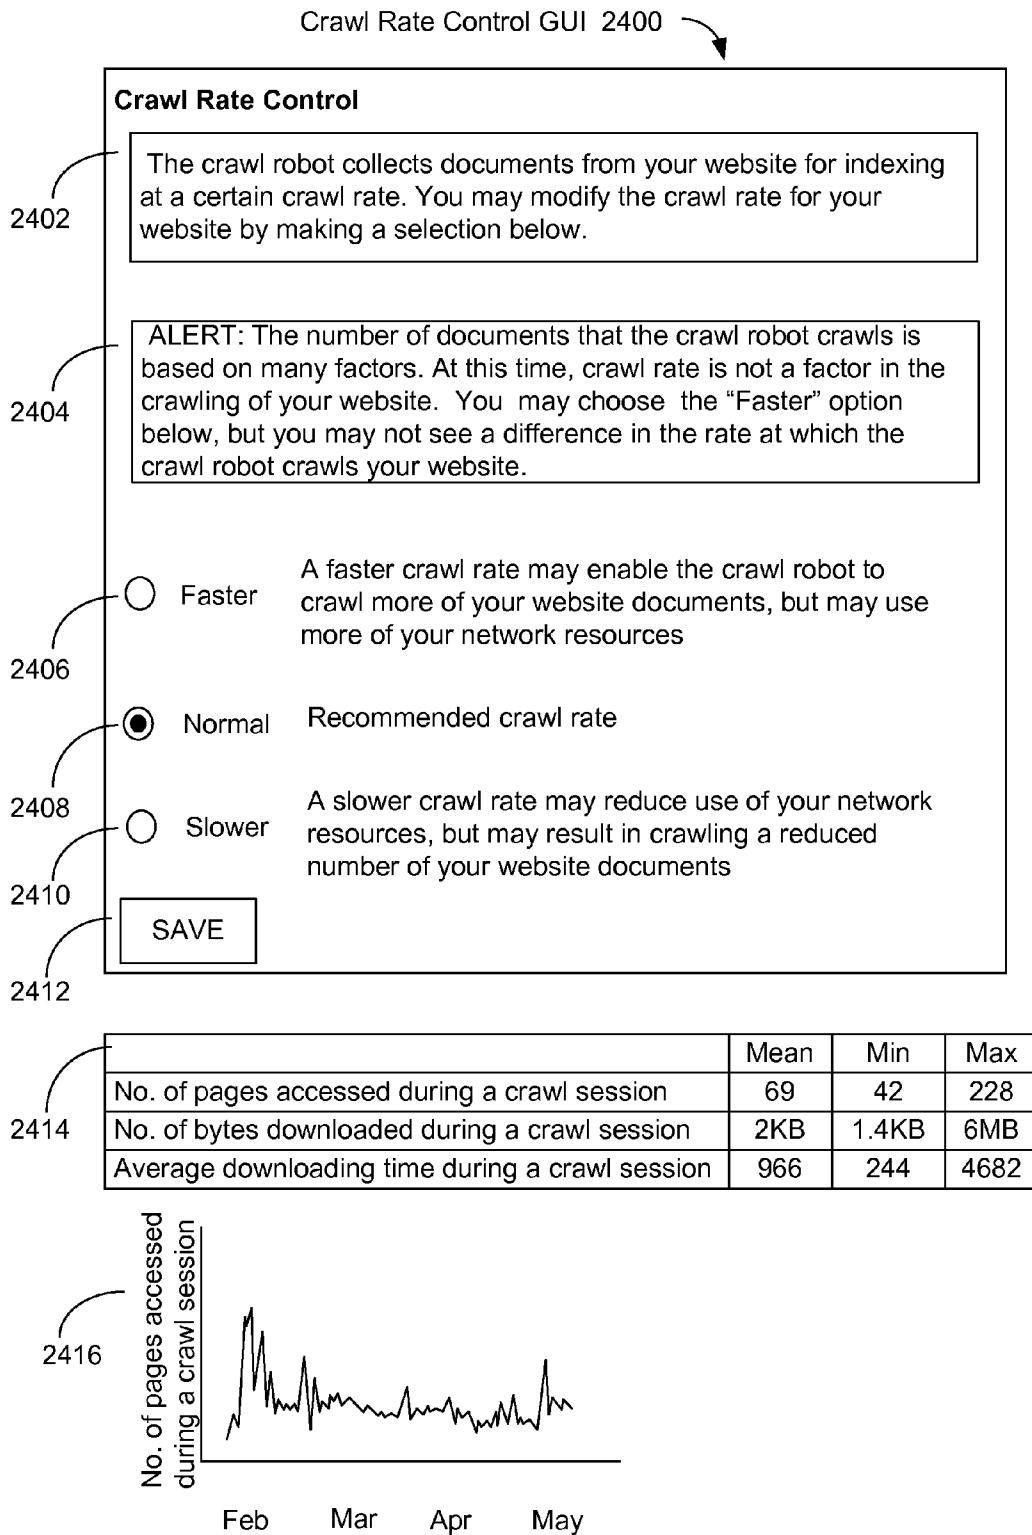


Figure 24



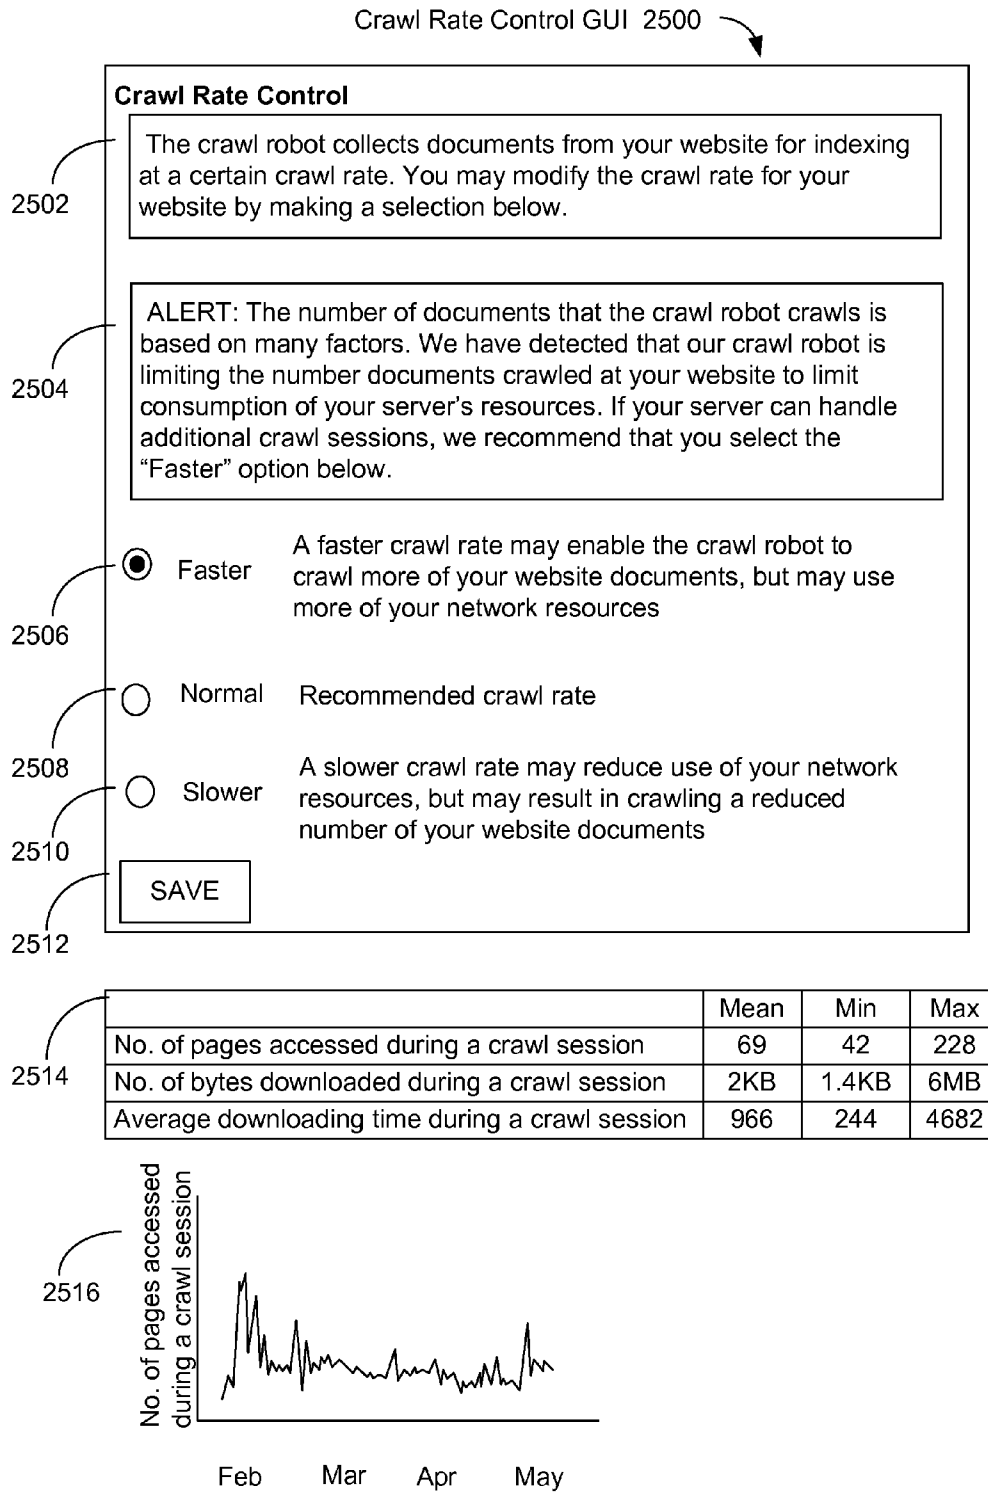


Figure 25

# SYSTEM AND METHOD FOR ENABLING WEBSITE OWNER TO MANAGE CRAWL RATE IN A WEBSITE INDEXING SYSTEM

## RELATED APPLICATIONS

This application is a continuation of U.S. patent application Ser. No. 12/554,757, filed Sep. 4, 2009, now U.S. Pat. No. 8,032,518, which is continuation of U.S. patent application Ser. No. 11/549,075, filed Oct. 12, 2006, now U.S. Pat. No. 7,599,920, which are hereby incorporated herein in their entireties.

This application incorporates by reference the following applications in their entirety: U.S. Provisional Patent Application No. 60/686,492, titled "Sitemap Generating Client for Web Crawler," filed May 31, 2005; U.S. patent application Ser. No. 11/172,764, titled "Web Crawler Scheduler that Utilizes Sitemaps from Websites," filed Jun. 30, 2005; and U.S. Provisional Patent Application No. 60/686,494, titled "Web Crawler Scheduler that Utilizes Sitemaps from Websites," filed May 31, 2005.

## TECHNICAL FIELD

The disclosed embodiments relate generally to search engines. More particularly, the disclosed embodiments relate to a management of crawl rate with respect to a website in a website indexing system.

## BACKGROUND

A search engine is a tool that identifies documents, typically stored on hosts distributed over a network, which satisfy search queries specified by users. Web search engines work by storing information about a large number of documents (such as web pages) which they retrieve from the World Wide Web (WWW) via a web crawler. Crawling a website during high traffic periods and/or excessive load during crawling can deplete network resources from the website, rendering the website less accessible to others. On the other hand, it may be that not enough crawling is performed for the website. Thus, it may be advantageous to enable a website owner to control the rate at which the website is crawled.

## BRIEF DESCRIPTION OF THE DRAWINGS

FIG. 1 is a conceptual diagram of a process for generating a sitemap for a website, according to some embodiments of the invention.

FIG. 2 is a block diagram illustrating a website server, according to some embodiments of the invention.

FIG. 3 is a block diagram illustrating a data structure for storing sitemap generator control parameters, according to some embodiments of the invention.

FIG. 4 is a flow chart illustrating a process for generating a sitemap, according to some embodiments of the invention.

FIG. 5 is a flow chart illustrating another process for generating a sitemap, according to some embodiments of the invention.

FIG. 6 is a flowchart illustrating a process for generating a differential sitemap, according to some embodiments of the invention.

FIG. 7 is a block diagram illustrating a web crawler system, according to some embodiments of the invention.

FIG. 8 is a block diagram illustrating a data structure in a domain name database, according to some embodiments of the invention.

FIG. 9 is a block diagram illustrating an alternate data structure in a domain name database, according to some embodiments of the invention.

FIG. 10 is a flowchart illustrating a process for using a preferred domain name when indexing documents, according to some embodiments of the invention.

FIG. 11 is a flowchart illustrating a process for enabling a website owner to select a preferred domain name, according to some embodiments of the invention.

FIGS. 12, 13, and 14 are schematic illustrations of a GUI for managing domain names, according to some embodiments of the invention.

FIG. 15 is a block diagram illustrating a data structure in a registered owner database, according to some embodiments of the invention.

FIG. 16 is a block diagram illustrating an alternate data structure in a site owner verification database, according to some embodiments of the invention.

FIGS. 17 and 18 are flowcharts illustrating a process for verifying website ownership, according to some embodiments of the invention.

FIGS. 19 and 20 are schematic illustrations of a GUI for managing site ownership verification, according to some embodiments of the invention.

FIG. 21 is a block diagram illustrating a data structure in a crawl rate database, according to some embodiments of the invention.

FIGS. 22 and 23 are flowcharts illustrating a process for controlling crawl rate, according to some embodiments of the invention.

FIGS. 24 and 25 are schematic illustrations of a GUI for controlling crawl rate, according to some embodiments of the invention.

Like reference numerals refer to corresponding parts throughout the drawings.

## DESCRIPTION OF EMBODIMENTS

A web server of a website generates a sitemap of URLs (Universal Resource Locators) or URIs (Universal Resource Identifiers) that may be crawled by a web crawler. The sitemap may also include metadata associated with the URLs and/or URIs included in the sitemaps. The web server sends a notification to a remote computer associated with a web crawler. In this context the web server may be considered to be the client and the web crawler is considered to be the server in a client-server relationship. The remote computer accesses the sitemap and uses the information included within the sitemap to identify documents stored at the web server to include in a crawl.

In context of search engines and network crawlers, the term "crawl" means to download and process documents so as to build an index or other listing of documents. "A crawl" is an iteration of a crawl process, during which a set of documents are downloaded and processed. The term "crawl" sometimes also means identifying and processing links in the downloaded documents so as to build a link map or other representation of links between the crawled documents. Typically, the purpose of a crawl is to compile the information in a set of documents for a search engine, thereby enabling the search engine to identify documents satisfying search queries.

FIG. 1 is a block diagram illustrating a website 100, according to some embodiments of the invention. The website 100 includes a website file system 102, sitemap generator control parameters 104, a sitemap generator 106, a sitemap update module 108, a sitemap notification module 110, sitemaps 114, and sitemap indexes 112. The file system 102 may be implemented in some embodiments using any of a number of file systems, including distributed file systems in which files are stored on multiple computers. In other

3

embodiments, the file system **102** may be implemented using a database or search engine that produces documents in response to queries. It should be appreciated, however, that FIG. 1 is intended more as functional description of the various features which may be present in a website than as a structural schematic of the embodiments described herein. In practice, and as recognized by those of ordinary skill in the art, items shown separately could be combined and some items could be separated. For example, some items shown separately in FIG. 1 could be implemented on single servers and single items could be implemented by one or more servers. As recognized by those of ordinary skill in the art, a website may be implemented on a single server, such as a web server, or on a plurality of servers, such as a plurality of web servers. The actual number of servers used to implement a website server and how features are allocated among them will vary from one implementation to another, and may depend in part on the amount of data traffic that the system must handle during peak usage periods as well as during average usage periods. For convenience of explanation, websites will be described below as if they are implemented on a single web server.

The website file system **102** organizes the documents that are stored at the web server. A document stored at the website may be any machine-readable file that includes text, graphics, video, audio, etc., or any combination thereof. Examples of documents that may be stored at a website include, without limitation, web pages, images, video files, audio files, Portable Document Format (PDF) files, plain text files, executable files, presentation files, spreadsheets, word processor documents, and so forth.

The documents stored at the website **100** may be organized in a hierarchical structure. That is, the documents are organized into a tree of nested directories, folders, or paths (hereinafter the "directory tree"). The directory tree includes a root directory/folder/path, and the root may have subdirectories/subfolders/subpaths nested within. The subdirectories/subfolders/subpaths may also have further subdirectories/subfolders/subpaths nested within, thus forming a directory tree. Each document is stored in a directory/folder/path in the directory tree. Each directory/folder/path and each document is a node in the tree. The file system also stores metadata associated with the documents, such as a last modification date, a last access date, document permissions, and the like. In some embodiments, the file system may also include a database of documents and associated metadata. Documents in the database may be accessed by executing a query of the database as well as, or instead of, traversing the directory tree.

Each document stored at the website may be identified and/or located by a locator. In some embodiments, the locator is the Uniform Resource Locator (URL) of the document. In some other documents, alternative manners of identification (e.g., URIs) or addressing may be used. The URL of a document may be derived from its location in the file system. The URL of a document may be based on the directory/folder/path, or the location in a database, or on the query used to retrieve the document from the database in which the document is stored. That is, each document in a directory/folder/path or a database location may be mapped to an URL. In some embodiments, the URLs may be used by computers external to the website, such as remote computers associated with web crawlers, to access the documents in the file system that are open to external access. For convenience of explanation, the document locators are described below as if they are URLs.

The sitemap generator **106** generates sitemaps and, optionally, one or more sitemap indexes of the website. The

4

sitemaps, further details of which are described below, may be used by web crawlers to schedule its crawls of documents stored at the web server. Sitemap indexes, further details of which are described below, encapsulate one or more sitemaps. The sitemap index may contain a list of sitemaps. The sitemap generator **106** generates sitemaps by accessing one or more sources of document information. In some embodiments, the sources of document information include the file system **102**, access logs, pre-made URL lists, and content management systems. The sitemap generator may gather document information by simply accessing the website file system **102** and collecting information about any document found in the file system **102**. For instance, the document information may be obtained from a directory structure that identifies all the files in the file system. The sitemap generator **106** may also gather document information by accessing the access logs (not shown) of the website. The access logs record accesses of documents by external computers. An access log may include the URLs of the accessed documents, identifiers of the computers accessing the documents, and the dates and times of the accesses. The sitemap generator **106** may also gather document information by accessing pre-made URL lists (not shown). The pre-made URL lists list URLs of documents that the website operator wishes to be crawled by web crawlers. The URL lists may be made by the website operator using the same format as that used for sitemaps, as described below.

If the documents in the website are managed via a content management system, the sitemap generator **106** may gather document information by interfacing with the content management system and accessing the information stored within the content management system. Content management systems are well known in the art and need not be further described.

The sitemap generator control parameters **104** include predefined parameters that control the generation of sitemaps. Further information regarding the sitemap generator control parameters **104** is described below, in relation to FIG. 3.

The sitemap generator **106** generates sitemaps **114** and possibly one or more sitemap indexes **112**. The sitemap **114** and sitemap index **112** may be generated using any suitable format and language. In some embodiments, the sitemap is generated in Extensible Markup Language (XML) format, using predefined XML tags. In some other embodiments, other formats may be used, including but not limited to plain text files, comma-separated value files, and semicolon-separated value files. For convenience of description, sitemaps and sitemap indexes below are described as formatted using XML.

A sitemap index **112** is a document that lists one or more sitemaps **114**. The sitemap generator **106**, when generating sitemaps for the website, may generate multiple sitemaps, with each sitemap listing a subset of URLs of documents that may be crawled, rather than listing URLs of all documents that may be crawled in one sitemap. In such a situation, the sitemap generator **106** may also generate a sitemap index **112** to list the multiple sitemaps and their URLs. A sitemap index may include starting and ending tags (e.g., XML tags such as <sitemapindex> and </sitemapindex>, not shown in the Figures) defining the beginning and end of the sitemap index **112**. The sitemap index also includes URLs of each sitemap listed in the sitemap index. The sitemap index may also include optional metadata for respective sitemap URLs in the sitemap index. For instance, the metadata may include a last modification date of the respective sitemap. Each sitemap URL and any respective associated metadata are enclosed by starting

and ending tags that define the beginning and end of a sitemap record **114** in the sitemap index **112**.

In addition to the list of sitemaps, in some embodiments a sitemap index may optionally contain a list of site-specific information **140** (also called “per-site information”) that applies to an entire website. For example, a sitemap index may contain a list of time intervals and a rate at which the crawler should crawl the respective website (e.g.,

---

```
<crawl_rate from=08:00UTC to=17:00UTC>medium</crawl_rate>
<crawl_rate from=17:00UTC to=8:00UTC>fast</crawl_rate>.
```

---

In other examples, a sitemap index contains geographic information identifying a geographic location associated with the website (e.g., <location>latitude, longitude</location>), and/or it may contain language information identifying one or more languages supported by or otherwise associated with the respective website (e.g., <language>German</language>). In some embodiments, per-site information may also be present in the sitemaps references in a sitemap index file. If both the sitemap index and a referenced sitemap contain per-site information for the same property (e.g., crawl rate), then the value specified in the sitemap overrides the value specified in the sitemap index. In other embodiments, per-site information may be specified in a sitemap index or sitemap using syntax other than the examples given here.

In one embodiment, the sitemap generator **106** of a website generates a new sitemap at regular intervals, such as daily or weekly. Each new sitemap generated, after a first (starting point) sitemap, only lists URLs that are new or modified since the prior sitemap was generated (i.e., have a creation date or modification date after the date that the last sitemap was generated). The term “date” is used here to mean date and time, and may be represented by a timestamp, such as an ISO 8601 compatible timestamp using UTC (coordinated universal time). In these embodiments, the sitemap index for the website lists all the sitemaps generated for the website. Optionally, a new starting point sitemap may be generated by the sitemap generator at larger intervals (e.g., weekly or monthly) than the intervals at which the update sitemaps are generated. Each time a new sitemap is generated and added to the sitemap index, a notification is sent to one or more search engines or crawlers.

A sitemap **114** is a document that lists the URLs of documents in a website that may be crawled by a web crawler. A sitemap **114** includes a list of URLs and, optionally, additional information, such as metadata, for respective listed URLs. A sitemap **114** may include starting and ending tags **116** that define the beginning and end of the sitemap. The sitemap includes one or more URL records **118**. The beginning and end of each URL record **118** are defined by a starting tag **120** and an ending tag **130**. Each URL record **118** includes the URL **122** of a document that may be crawled. A URL record **118** may also include optional metadata associated with the respective URL. The optional metadata may include one or more of the following: a last modification date **124** of the document specified by the URL, a change frequency **126** (also called the update rate) of the document specified by the URL, a document title **127**, a document author **129**, and a priority **128** of the document specified by the URL. The change frequency **126** and priority **128** may be specified by the website operator.

The change frequency **126** is a descriptor of how often a document’s content is expected to change. The descriptor is one of a predefined set of valid descriptors. In some embodi-

ments, the set of change frequency descriptors includes “always,” “hourly,” “daily,” “weekly,” “monthly,” “yearly,” and “never.” The change frequency **126** provides a hint to the crawler as to how often the document changes and the crawler can use the hint to schedule crawls of the document accordingly. A crawler, however, may crawl the document in a way that is inconsistent with the specified change frequency. For example, a crawler may crawl a document marked as “hourly” less frequently than a document marked as “yearly.” The actual crawl frequency of a document may be based on the document’s importance (as represented by a score, such as PageRank), actual observed changes in the document (or lack thereof) as observed by the crawler, and other factors, as well as the change frequency specified in a sitemap.

The priority **128** is a value that specifies the relative priority of the document identified by the URL **122**. The priority **128** may be the priority relative to other documents listed in the same sitemap **114**, relative to other documents stored in the same web server as the document, or relative to all documents in the website. In some embodiments, the range of priority values is 0.0 to 1.0, inclusive, with 0.5 as the default value, 0.0 as the lowest relative priority, and 1.0 as the highest relative priority. In some other embodiments, other priority scales, such as 0 to 10, may be used. The priority may be used by a crawler to determine which documents in the website should be crawled first. The crawler may ignore or modify the priority values in a sitemap when those priority values fail to meet predefined criteria (e.g., a requirement that the priority values in a sitemap or set of sitemaps for a website have a predefined average value, such as 0.5). In some embodiments, the priority may also be used when indexing documents.

The sitemap generator **106** may also interact with a sitemap update module **108** and a sitemap notification module **110**. The sitemap notification module **110** sends a notification to a remote computer associated with a web crawler whenever a new or updated sitemap is available at the website. The notification includes the URL of the sitemap, so that the remote computer can access the sitemap. If the website uses sitemap indexes, the notification, in some embodiments, includes just the URL of the sitemap index. The remote computer accesses the sitemap index and identifies the URLs of the sitemaps from the sitemap index. In some other embodiments, the notification may include the actual sitemap index, in which case the remote computer need not access the sitemap index at the website.

The sitemap update module **108** may generate a differential sitemap based on a difference between a previously generated sitemap and a current sitemap. Further information regarding differential sitemaps is described below, in relation to FIG. 6.

FIG. 2 is a block diagram illustrating a website server **200**, according to some embodiments of the invention. The website server **200** (or “web server”) typically includes one or more processing units (CPU’s) **202**, one or more network or other communications interfaces **204**, memory **210**, and one or more communication buses or signal lines **212** for interconnecting these components. The website server **200** optionally may include a user interface **205**, which may comprise a display device, mouse, and/or a keyboard. The memory **210** includes high-speed random access memory, such as DRAM, SRAM, DDR RAM or other random access solid state memory devices; and may include non-volatile memory, such as one or more magnetic disk storage devices, optical disk storage devices, flash memory devices, or other non-volatile solid state storage devices. Memory **210** may optionally include one or more storage devices remotely located from the CPU(s) **202** (e.g., network attached storage). In some

7

embodiments, the memory **210** stores the following programs, modules and data structures, or a subset thereof:

an operating system **214** that includes procedures for handling various basic system services and for performing hardware dependent tasks;

a network communication module **216** that is used for connecting the website server **200** to other computers via the one or more communication network interfaces **204** and one or more communication networks, such as the Internet, other wide area networks, local area networks, metropolitan area networks, and so on;

a sitemap generation module **106** that generates sitemaps; sitemap control parameters **104** that control or guide sitemap generation;

a sitemap index **112** that lists URLs of sitemaps stored at the website server **200**;

one or more sitemap(s) **114** that list URLs of documents that may be crawled;

a website file system **102** that stores and organizes documents;

a domain name database **222**;

a registered owner database **224**;

a site owner verification database **226**, discussed below with reference to FIGS. **15-16**;

a crawl rate database **228**, discussed below with reference to FIG. **21**;

a domain name module **230**;

a site owner verification module **232**, discussed below with reference to FIGS. **17-20**; and

a crawl rate control module **234**, discussed below with reference to FIGS. **21-25**.

Each of the above identified elements may be stored in one or more of the previously mentioned memory devices, and corresponds to a set of instructions for performing a function described above. The above identified modules or programs (i.e., sets of instructions) need not be implemented as separate software programs, procedures or modules, and thus various subsets of these modules may be combined or otherwise re-arranged in various embodiments. In some embodiments, memory **210** may store a subset of the modules and data structures identified above. Furthermore, memory **210** may store additional modules and data structures not described above.

Although FIG. **2** shows a “website server,” FIG. **2** is intended more as functional description of the various features which may be present in a set of servers than as a structural schematic of the embodiments described herein. In practice, and as recognized by those of ordinary skill in the art, items shown separately could be combined and some items could be separated. For example, some items shown separately in FIG. **2** could be implemented on single servers and single items could be implemented by one or more servers. The actual number of servers used to implement a website server and how features are allocated among them will vary from one implementation to another, and may depend in part on the amount of data traffic that the system must handle during peak usage periods as well as during average usage periods.

FIG. **3** is a block diagram illustrating a data structure for storing sitemap generator control parameters, according to some embodiments of the invention. The sitemap generator control parameters **104** control the generation of sitemaps and sitemap indexes. Each of the parameters may be specified by the operator of the website. The parameters may include one or more of the following:

8

one or more sitemap base URLs **302**, which specifies the location(s) from which the sitemaps may be accessed by remote computers associated with web crawlers;

file path to URL mapping(s) **304**, which map directories/paths/folders or database locations in the file system **102** to externally accessible URLs (an exemplary path to URL mapping is P:/A/B/\*.\*> www.website.com/qu/\*.\*);

URL exclusion pattern(s) **306**, which specify classes of URLs that are to be excluded from inclusion in sitemaps (e.g., an exclusion pattern of “www.website.com/wa/\*.pr1” would indicate that all “pr1” files in the “/wa” portion of www.website.com are to be excluded from the sitemap(s));

URL pattern(s) with update rates **308**, which specify classes of URLs and update rates (change frequencies) for respective URL classes (e.g., www.website.com/qu/a\*.pdf> daily would indicate that the files meeting the specified patterns are expected to be updated daily.

notify URL(s) **310**, which specify the URLs of remote computers associated with web crawlers to which new sitemap notifications may be transmitted;

pointer(s) to URL list(s) **312**, which point to pre-made URL lists;

pointer(s) to URL access log(s) **314**, which point to URL access logs;

pointer(s) to directory or directories **316**, which point to directories/folders/paths or database locations in the file system **102**; and

(optional) preferred crawl time(s) **318**, which specifies preferred times of the day for web crawlers to crawl the website. In some embodiments, blackout periods during which the web crawler should not crawl the website may be specified.

It should be appreciated that the parameters listed above are merely exemplary and that additional and/or alternative parameters may be included.

FIG. **4** is a flow chart illustrating a process for generating a sitemap, according to some embodiments of the invention. As described above, a source of information about documents stored on the website is the access logs of the website. The access logs of the website are accessed (**402**). The access logs may be found by following pointers **314** to URL access logs. The access logs are scanned for non-error URLs (**404**). Non-error URLs are URLs that properly specify an existing and accessible document. Thus, for example, a URL for a document that is no longer on the website may be deemed an error URL. A list of URLs is generated (**406**). The list includes the non-error URLs found in the access logs. The list optionally may also include document popularity information derived from the access logs. The document popularity information may be determined based on the numbers of accesses each non-error URL has. The document popularity information serves as an additional hint of which documents are to be given a higher priority during crawling (e.g., scheduled to be crawled first, or more likely to be crawled than lower priority documents), based on which documents are in high demand (i.e., are accessed more often).

After the list of URLs is generated, the list may optionally be filtered for excluded URLs (**410**). The URL exclusion patterns **406** from the sitemap generator control parameters **104** may be used as the filter that is applied to the list of URLs (**408**). Alternately, URL exclusion patterns may be obtained elsewhere, or permanently encoded in a customized sitemap generator for the website. URLs in the list that match any of the URL exclusion patterns are removed from the list.

Optional update rate information may be added to the list of URLs, for the respective URLs in the list (412). In some embodiments, the update rates may be obtained from the sitemap generator control parameters 104, or more particularly, the URL pattern(s) with update rates 308 (408).

Last modification dates and times for the respective URLs in the list of URLs are added (414). The last modification dates may be obtained from the file system, which may be a database and/or a directory tree 416, as described above.

In an alternate embodiment, a sitemap strategy object 415 controls the filtering operation 410, update rate information adding operation 412 and the last modification date adding operation 414, using information obtained from a database 416 and/or the sitemap generator control parameters 408. In some embodiments, the sitemap strategy object determines which URLs (or URIs) to filter and which attributes to add to specific URLs (or URIs) by performing a database query on the underlying database 416.

The sitemap is generated from the resulting list of URLs, including any last modification date/time information, optional popularity information, and optional update rate information that has been included or obtained for the listed URLs (418). Within the sitemap, the metadata for the URLs listed in the sitemap comprises the last modification date/time information, optional popularity information, and optional update rate.

FIG. 5 is a flow chart illustrating another process for generating a sitemap, according to some embodiments of the invention. The process of FIG. 5 is similar to that of FIG. 4, with the difference being that, in the process of FIG. 5, the initial source of document information is the file system database or directory tree (502), rather than access logs. A scan of the database or a traversal of the directory tree is performed (504). From the database scan or directory tree traversal, a list of URLs and associated last modification dates are obtained (506). The list may be optionally filtered for excluded URLs (508), using the URL exclusion patterns from the sitemap generator control parameters as the filters (512). Optionally, additional metadata, such as document update rate information associated with the respective URLs in the list of URLs may be added (510). The update rate information may be taken from the sitemap generator control parameters (512). The sitemap is generated from the list of non-excluded URLs, the last modification date information, and additional information such as the update rate information.

In an alternate embodiment, a sitemap strategy object 515 controls the filtering operation 508 and/or adding metadata 510 to the list of URLs or URIs in the sitemap 514 using information obtained from the underlying database 502 and/or the sitemap generator control parameters 512. In some embodiments, the sitemap strategy object 515 determines which URLs (or URIs) to filter and which attributes to add to specific URLs (or URIs) by performing a database query on the underlying database 502.

It should be appreciated that the sitemap generation processes illustrated in FIGS. 4 and 5 may be adapted to use alternative sources of document information and/or use multiple sources of document information. For example, the sitemap generator may initially extract URLs from one or more pre-made URL lists or from a content management system associated with the website. Regardless of the source from which the URLs are extracted, the sitemap generator may collect document metadata from as many of the sources of document information as needed. For example, the sitemap generator may extract URLs from pre-made URL lists, obtain the last modification dates from the file system and obtain document popularity information from the access logs. Any

suitable combination of sources of document information may be used to generate the sitemaps.

FIG. 6 is a flowchart illustrating a process for generating a differential sitemap, according to some embodiments of the invention. A differential sitemap is a sitemap that is generated based on a difference between a previously generated sitemap and a current sitemap. The differential sitemap includes URLs that were not included in the previously generated sitemap and URLs that were included in the previously generated sitemap but which have new or updated metadata. For example, a URL with an updated last modification date would be included in the differential sitemap. The presence of an updated last modification date for the URL means that the document at the respective URL has been updated since the previously generated sitemap.

A current sitemap (602) and a previously generated sitemap (604) are processed by a differential sitemap generator, such as the sitemap update module 108 (606). The differences between the two sitemaps are ascertained and a differential sitemap is generated (608).

The sitemaps, if used by web crawlers, may lead to greater crawl coverage, since the sitemaps may include documents, such as documents that are accessible only by a database query, that are not accessible by following links. The sitemaps may also provide last modification dates. Web crawlers may use the last modification dates to determine if a document has changed, and may thereby avoid crawling documents whose contents have not changed. The use of sitemaps to avoid crawling unchanged documents can make web crawlers and network crawlers significantly more efficient. The sitemaps also include information from which web crawlers may determine which documents to crawl first and how much load to put on the web server during crawling. This, too, may help conserve network resources.

FIG. 7 is a block diagram illustrating a web crawler system 700, according to some embodiments of the invention. The web crawler system 700, which may be a part of and/or associated with a search engine, crawls locations corresponding to documents stored in web servers. As discussed above, a document stored at the website may be any machine-readable file that includes text, graphics, video, audio, etc., or any combination thereof.

The sitemap crawler 705 accesses sitemaps generated by websites or web servers. The sitemap crawler 705 receives sitemap notifications. The sitemap notifications 730 are received from web servers or websites that have documents available for crawling. A notification from a web server or website informs the sitemap crawler that one or more sitemaps listing URLs of documents that may be crawled are available for access. The notification may include the URL of a sitemap, or the URLs of two or more sitemaps. The notification may include the URL of a sitemap index, or it may include the content of a sitemap index. In some embodiments, the notification may include the sitemap index (112, FIG. 1) or the full sitemap. The sitemap crawler 705 may access the sitemap index at the sitemap index URL to learn the URLs of the sitemaps and then access the sitemaps.

The sitemap crawler 705 accesses sitemaps from web servers or websites and may store copies of the accessed sitemaps in a sitemap database 732. The sitemap database 732 stores the sitemaps and information associated with the sitemaps, such as the web servers and/or websites with which the sitemaps are associated, the last modification dates of the sitemaps, and update rate information associated with the sitemaps.

Accessed sitemaps are provided to a sitemap processing module 734 for processing. The sitemap processing module

**734** processes the sitemaps and identifies the URLs and associated metadata **736**. The sitemaps are a source of URLs and associated metadata information for the URL scheduler **702**. In some embodiments, an optional, additional source of URLs and associated metadata is direct submission **703** by users.

The URL scheduler **702** determines which URLs will be crawled in each crawling session. The URL scheduler **702** may store that information in one or more data structures (not shown), such as a set of list data structures. In some embodiments, the URL scheduler **702** allocates URLs to segments of the data structure, where the segments correspond to the crawl sessions. In these embodiments, the URL scheduler **702** also determines which URLs within each segment are to be crawled. In some embodiments, there may be a number of URL schedulers **702**, which are run prior to each segment being crawled. Each scheduler **702** is coupled to a corresponding URL manager **704**, which is responsible for managing the distribution of URLs to URL server **706**. Alternatively, each URL scheduler **702** may be coupled to two or more of the URL managers, enabling the URL distribution function for each crawl session to be spread over multiple URL managers. The URL schedulers **702** are adapted to receive URLs and metadata **736** extracted from sitemaps.

A controller **701** selects a segment for crawling. The selected segment is referred to hereinafter as the “active segment.” Typically, at the start of each session, the controller **701** selects a different segment as the active segment so that, over the course of several sessions, all the segments are selected for crawling in a round-robin manner.

A query-independent score (also called a document score) is computed for each URL by URL page rankers **722**. The page rankers **722** compute a page importance score (sometimes called a document importance score) for a given URL. In some embodiments, the page importance score is computed by considering not only the number of URLs that reference a given URL but also the page importance score of such referencing URLs. Page importance score data is provided to URL managers **704**, which pass a page importance score for each URL to URL server **706**, robots **708**, and content processing servers **710**. One example of a page importance score is PageRank, which is used the page importance metric used in the Google search engine. An explanation of the computation of PageRank is found in U.S. Pat. No. 6,285,999, which is incorporated by reference herein in its entirety as background information. In some embodiments, information from the sitemaps may be incorporated into the computation of the page importance score. An example of sitemap information that may be incorporated into the page importance score is the priority **128**.

From time to time, the URL server **706** requests URLs from the URL managers **704**. In response, the URL managers **704** provide the URL server **706** with URLs obtained from the data structure. The URL server **706** then distributes URLs from the URL managers **704** to crawlers **708** (hereinafter also called “robots” or “bots”) to be crawled. A robot **708** is a server that retrieves documents at the URLs provided by the URL server **706**. The robots **708** use various known protocols to download pages associated with URLs (e.g., HTTP, HTTPS, Gopher, FTP, etc.). According to certain embodiments, a crawl rate control module **742** sends crawl rate information for a specified website to the per-site information database **740**. In some embodiments, a robot **708** retrieves, from the per-site information database **740** crawl rate and/or crawl interval information for a specified website, and then uses the retrieved information to control the rate at which pages associated with the URLs are accessed or downloaded

by robot **708** from the website. In other embodiments, the URL server **706** retrieves the crawl rate information for a specified website from the per-site information database **740** and distributes the crawl rate information to one or more robots. According to certain embodiments, crawl rate may be controlled by a site owner as described in greater detail herein with reference to FIGS. **21-25**.

Pages obtained from URLs that have been crawled by robots **708** are delivered to content processing servers **710**, which perform a number of tasks. In some embodiments, these tasks include indexing the content of the pages, generating records of the outbound links in the pages, detecting duplicate pages, and creating various log records to record information about the crawled pages. In one embodiment, these log records are stored in log files, including link logs **714**, status logs **712**, and other logs **716**. The link logs **714** include a link record for each document obtained from a URL by a robot **708** and passed to the content processing servers **710**. Each link log **714** record identifies all the links (e.g., URLs, also called outbound links) that are found in the document associated with the record and the text that surrounds the link. The information in the link logs **714** is used by the content processing servers **710** to create link maps **720**. The records in the link map **720** are similar to records in the link logs **714** with the exception that text is stripped and the records are keyed by a “fingerprint” of the normalized value of the source URL. In some embodiments, a URL fingerprint is a 64-bit integer determined by applying a hash function or other one way function to a URL. The bit-length of the URL fingerprint may be longer or shorter than 64 bits in other embodiments. The records in each link map **720** may optionally be sorted or keyed by a fingerprint. The link maps **720** are used by the page rankers **722** to compute or adjust the page importance score of URLs. In some embodiments, such page importance scores may persist between sessions.

The status logs **712** log the status of the document processing performed by the content processing servers **710**. The status logs may include URL status information **728** (e.g., whether a document existed at the specified URL, last modification date information, and update rate information). The URL status information may be transmitted to the URL scheduler(s) **702**. The URL scheduler(s) may use the URL status information to schedule documents for crawling.

In some embodiments, the content processing servers **710** also create anchor maps **718**. The anchor maps **718** maps the “anchor text” in hyperlinks to the URL of the target URLs of the hyperlinks. In documents that use HTML tags to implement the hyperlinks, the anchor text is the text located between a pair of anchor tags. For example, the anchor text in the following pair of anchor tags is “Picture of Mount Everest”:

```
<A href="http://www.website.com/wa/me.jpg"> Picture
of Mount Everest</A>.
```

In some embodiments, sitemap-supplied document metadata may also be used for creating anchor maps. For example, document metadata such as the document title **127** (FIG. 1), document author **129**, or document description (not shown) may be used to create the anchor maps. However, it should be appreciated that any field appearing in a sitemap may be included in an anchor map.

In some embodiments, the records in an anchor map **718** are keyed by the fingerprints of outbound URLs present in link log **714**. Thus, each record in an anchor map **718** comprises the fingerprint of an outbound URL and the anchor text that corresponds to the URL in link log **714**. Anchor maps **718** are used by the indexer(s) **724** to facilitate the indexing of “anchor text” as well as to facilitate the indexing of URLs that

13

do not contain words. For example, consider the case in which the target document at an outbound URL (e.g., the URL in the above example) is a picture of Mount Everest and there are no words in the target document. However, anchor text associated with the URL, "Picture of Mount Everest" is included in an index **726**, sometimes called an inverse document index, thereby making the target document accessible via a search engine using the index **726**.

The anchor maps **718** and other logs **716** are transmitted to indexer(s) **724**. The indexer(s) **724** use the anchor maps **718** and other logs **716** to generate index(es) **726**. The index(es) are used by the search engine to identify documents matching queries entered by users of the search engine.

FIG. **8** is a block diagram illustrating a domain name database **800** according to some embodiments of the invention. Domain name database **800** includes one or more website records **802**, each of which stores domain name information (and possibly other information as well) for a respective website. All the verified domain names associated with a respective website are considered to be equivalent domain names in that any of the verified domain names can be used as the domain name portion of a webpage address or URL for addressing any webpage of the website. The website owner or manager, however, may prefer that a search engine index all pages of the website using a single, preferred domain name, even though links (to webpages in the website) in various webpages or other documents may use other ones of the website's verified domain names.

A respective website record **802** includes website information **804** and one or more domain name preference records **806**. Website information **804** includes a preferred domain name **808** and optionally, additional information. A respective domain name preference record **806** includes a domain name mapping **810**, a verification status **812**, a list of verifying owners **814**, and optionally, additional information. Preferred domain name **808** is the current preferred domain name for the website corresponding to the website record **802** in which the preferred domain name is stored. Each domain name mapping **810** identifies or specifies a respective domain name that is mapped to the preferred domain name. The list of verifying owners **814** identifies the one or more verified owners who have provided or confirmed the mapping shown in domain name mapping **810**. As a non-limiting example, assume that A, B, C and X are equivalent domain names for a website. Further assume, the selected preferred domain name is X. Domain name mapping **810** of domain name preference record **806-1** may show the mapping of A→X. Similarly, domain name mapping **810** of domain name preference record **806-2** may show the mapping of B→X, etc. If at least one verifying owner from the list of verifying owners **814** is currently verified, then verification status **812** shows that the domain name mapping **810** is verified as valid. If no verifying owner from the list of verifying owners **814** is currently verified, then verification status **812** shows that the domain name mapping **810** is not verified and is invalid. Further, according to certain embodiments, a determination is made as to whether the domain name mapping that is provided by the one or more verified owners is an appropriate mapping. For example, it is determined whether the content at the respective websites associated with the equivalent domain names as specified by the mapping are identical. As another non-limiting example, a determination may be made as to whether a "301 redirect" or a "302 redirect" has been set up for web pages from one domain name to an equivalent domain name as defined in the domain name mapping. According to certain embodiments, the system that implements or maintains the domain name database **800** ensures that the domain name

14

mappings do not contain cycles and/or chains. For example, the mappings A→X, X→Y are not allowed to coexist. Mappings such as A→X, X→Y, Y→A are also not allowed to coexist. However, according to certain other embodiments, the domain name mappings can contain chains, but not cycles.

FIG. **9** is a block diagram illustrating an alternate data structure for a domain name database **900**, according to some embodiments of the invention. The domain name database **900** includes one or more website records **902**. Website record **902** includes website information **904** and one or more name records **906** that correspond to one or more equivalent domain names for a given website. A respective name record **906** includes a domain name or domain name identifier **908**, a preferred flag **910**, a verification status **912**, a list of verifying owners **914**, and optionally, additional information. If the preferred flag **910** is set for a given domain name, then that domain name is the preferred domain name for the website corresponding to the website record **902** in which the preferred flag **910** is stored. The list of verifying owners **914** shows the list of verified owners that have identified or confirmed the domain name **908** as the preferred domain name. If at least one verifying owner from the list of verifying owners **914** is currently verified, then verification status **912** shows that the domain name **908** is a currently verified domain name for the website. If no verifying owner from the list of verifying owners **914** is currently verified, then verification status **912** shows that the domain name **908** does not qualify as a verified domain name for the website.

According to certain embodiments, a website owner can select a preferred domain name for indexing documents of his or her website. For purposes of explanation, assume that a website has two URLs that point to the same document on the website. The two URLs may be www.domain.com/p1.html and domain.com/p1.html. A search engine may not assume that www.domain.com/p1.html and domain.com/p1.html point to the same document. Further, even if the search engine determines that www.domain.com/p1.html and domain.com/p1.html point to the same document, the search engine does not know which domain name is preferred by the website owner for indexing the document of the website. As another non-limiting example, a website may use several equivalent domain names in the URLs associated with the website. For example, domain1.com, domain2.com, and shortcutdomain.com may all be equivalent to domain.com. In other words, the URLs domain1.com/P1.html, domain2.com/P1.html, shortcutdomain.com/P1.html and domain.com/P1.html all point to the same document (document P1.html) in the website. In this example, document P1.html represents any document in the website. During indexing of documents of such a website, a search engine may index document P1 as if it were several documents. Thus, the ranking value (and/or other metadata) of document P1 is diluted. Further, if the search engine had knowledge that the URLs domain1.com/P1.html, domain2.com/P1.html, shortcutdomain.com/P1.html and domain.com/P1.html point to the same document, and that domain.com is the preferred domain name, the URLs domain1.com/P1.html, domain2.com/P1.html, shortcutdomain.com/P1.html can be rewritten to domain.com/P1.html before commencing crawling of the website. As a consequence, the search engine crawls fewer documents and significant bandwidth savings may be realized.

According to certain embodiments, a mechanism may be provided to enable a website owner to specify a preferred domain name. Some non-limiting examples of such a mechanism include a user interface, meta-tags and a file stored in the website and having a predefined name (e.g., robots.txt) to



## 15

enable the website owner to specify a preferred domain name. As another non-limiting example, a predetermined protocol may be used to specify the preferred domain name. As another feature of some embodiments, the website owner may change the preferred domain name at any time. According to certain embodiments, the user specifying the preferred domain name is verified as a valid website owner. Instructions to use a domain name selected by a verified website owner as the preferred domain name may receive more weight than similar instructions from a non-verified user, according to some embodiments.

FIG. 10 is a flowchart illustrating a process for using a preferred domain name when indexing documents, according to some embodiments of the invention. A plurality of domain names including a preferred domain name is associated with a given website (1002). For a document in the website that has a specified URL, the specified URL is rewritten to a new URL that includes the preferred domain name (1004). For example, during a web crawl by a search engine's web crawler, the URL in a link (found in a previously crawled document) is rewritten to a new URL, thereby replacing the domain name of the URL from a first domain name to the preferred domain name of a respective website. Metadata for the document is generated using the re-written URL (1006). A non-limiting example of metadata for the document is page rank (sometimes written as "PageRank"). In addition, the document is indexed using the rewritten URL (1008). In some embodiments, indexing the document produces entries in the inverse index of a search engine. The inverse index's entries corresponding to the indexed document reference the rewritten URL. As a result, when the search engine returns a search result that includes the document, the URL of the document in the search result is the rewritten URL.

FIG. 11 is a flowchart illustrating a process for enabling a website owner to select a preferred domain name, according to some embodiments of the invention. A user is provided with a list of verified domain names that are stored in a domain name data store that is associated with the website (1102). The user is also provided with a mechanism, such as a user interface, to specify the preferred domain name (1104). If the domain specified by the user is from the list of verified domain names (1106—Yes), then the specified domain name is accepted as the preferred domain name for rewriting URLs and metadata of documents of the website (1110). If the specified domain name is not from the list of verified domain names (1106—No), then it is determined if the user owns the domain name specified by the user as the preferred domain name (1108). If it is determined that the user owns the domain name specified by the user (1108—Yes), then the user specified domain name is accepted as the preferred domain name for rewriting URLs and for generating metadata for documents of the website. If, however, it is determined that the user does not own the domain name specified by the user (1108—No), then the user is asked to specify another domain name (1112) and the process may resume at operation 1106.

According to certain embodiments, if the specified domain name is not from the list of verified domain names, then verification is performed on the specified domain name. Upon successful verification, the specified domain name is added to the list of verified domain names, and thus can be subsequently selected as a preferred domain name by using the process illustrated in FIG. 11, for example.

According to certain embodiments, the requests for preferred domain names by multiple users of a given website are analyzed to ensure that the requested domain name mappings do not result in chains and/or cycles as described herein with reference to FIG. 8. According to certain other embodiments,

## 16

domain name mappings that result in chains are accepted by the system. For example, for the domain name mapping  $A \rightarrow X$ ,  $X \rightarrow Y$ , the preferred domain name is Y.

FIG. 12 illustrates a GUI 1200 for managing domain names for a specified website. According to certain embodiments, a graphical user interface such as GUI 1200 is presented to a user to show the user the equivalent domain names for a specified website and to enable the user to select a preferred domain name. GUI 1200 shows:

instructions 1202 for using GUI 1200;

selection button 1204 for selecting a first domain name as the preferred domain name. In this example, it is assumed that there are only two equivalent domain names for the specified website;

selection button 1206 for selecting a second domain name as the preferred domain name;

selection button 1208 for electing not to specify a preferred domain name; and

selection button 1210 to activate the user's selection using GUI 1200.

GUI 1200 only shows two equivalent domain names for a specified website. However, it is understood by one skilled in the art, that a specified website may have a two or more equivalent domain names.

FIG. 13 shows a GUI 1300 that also is associated with managing domain names for a specified website. Upon activating the user's selection of a preferred domain name as described with reference to FIG. 12 herein, GUI 1300 shows confirmation information 1302 of the preferred domain name selected by the user. GUI 1300 may optionally include additional information 1304 such as a date when the preferred domain name was set. Further, GUI 1300 provides instructions 1306 for clearing the current domain name preference by selecting button 1308 to clear the preference.

FIG. 14 illustrates another non-limiting example of a preferred domain name GUI 1400. GUI 1400 shows domain name information 1402 explaining the role of a preferred domain name in the context of indexing a specified website. GUI 1400 also shows instructions 1404 for selecting a preferred domain name. Further, GUI 1400 provides a pull down list 1406 that lists a plurality of domain names associated with the specified website. A user can select a preferred domain name from list 1406 and then select the "ok" button 1408 to activate the selection. According to certain embodiments, the preferred domain name specified by the user is stored in a domain name database, such as domain name database 800 as described herein. For example, the preferred domain name selected by the user using either GUI 1300 or GUI 1400 may be stored as preferred domain name 808 for a specified website record 802. Further, once the preferred domain name is selected, the respective domain name preference record 806 of domain name database 800 is modified to store the mapping information for mapping a respective equivalent domain name to the specified preferred domain name in domain name mapping 810.

According to certain embodiments, a user may be verified as a valid owner of the website. A specified website may have one or more verified owners. A verified owner of a website may select a preferred domain name for the website in a manner as previously described herein. Further, a verified owner may perform other functions such as controlling a crawl rate for the website as described in greater detail herein with reference to FIGS. 21 through 25. Verification of website ownership is described in greater detail with reference to FIGS. 15 through 20, herein.

For a respective user, the system stores information such as the various websites for which the user is verified as an owner,

17

the time when the verification occurred, and the manner by which the user was verified. FIG. 15 is a block diagram illustrating a registered owner database 1500, according to some embodiments of the invention. The registered owner database 1500 includes one or more user records 1502. According to certain embodiments, a respective user record 1502 includes user information 1504, and one or more site records 1506. In some embodiments, a respective site record 1506 includes the site URL 1508, a location 1510 of the verification tag and/or the verification file, a timestamp 1512 associated with the verification of the user, the verification status 1514 of the user, and optionally, additional information. Site URL 1508 is the URL of the website for which the user is verified as one of the owners.

According to certain embodiments, the system provides a verification tag and/or a verification file that website owners can store on their respective websites. If a verification tag is used, the verification tag may be stored, for example, in a header of the home page of the website. Alternately, the verification tag may be stored in the website in a file having a predefined name (e.g., robots.txt or any other suitable file name). If a verification file is used, the verification file may be stored in the website.

According to some embodiments, each verification tag and/or verification file is specific to a respective owner. For example, in some embodiments, the verification tag for a respective website owner may include a verification value that is generated by (or on behalf of) the search engine and that is unique to the website owner. In one example, the verification value is produced by applying a one-way hash function to a name of the website owner concatenated with a seed value or other value. By maintaining secrecy with respect to the one-way hash function and/or the seed value or other value, attempts to generate non-authentic verification tags may be thwarted. In embodiments that use a verification file for verifying a respective website owner, the aforementioned verification value may be used as the name (or as a portion of the name) of the verification file.

Further, in some embodiments, the content of a verification tag or the name of a verification file may be produced in a manner that is transparent to the user and yet still secure. For example, the content of the verification tag or the name of the verification file can be digitally signed in a secure manner, using any appropriate digital signature methodology. The use of digital signature enables the verification tag or verification file to be authenticated by other users, including the search engine. Alternately, the content of the verification tag or the name of the verification file can be encrypted using public-private key encryption, which also allows authentication of the verification tag or verification file.

According to some embodiments, for a specified website that the system is crawling, the system stores information such as a list of users that are verified as owners of the specified website, the time when the verification occurred, and the manner by which each user was verified for the specified website. FIG. 16 is a block diagram illustrating a site owner verification database 1600, according to some embodiments of the invention. The site owner verification database 1600 includes one or more site records 1602. A respective site record 1602 corresponds to a respective website that the system may crawl. In some embodiments, a site record 1602 includes site URL 1604, and one or more user records 1606. Site URL 1604 is the URL associated with the respective website. User record 1606 includes a user identification such as user name or user ID 1608, a location 1610 of the verification tag and/or the verification file, a timestamp 1612 asso-

18

ciated with the verification of the user, the verification status 1614 of the user, and optionally, additional information.

FIG. 17 is a flowchart illustrating a process for verifying ownership rights of a website, according to some embodiments of the invention. Owners-information that is associated with currently verified owners of the website is stored (1702). As a non-limiting example, owners-information may include a list of verification tags and/or verification files associated with each of the respective owners of the website. The list of verification tags and/or verification files may be stored either in verification tag/file location 1510 or 1610 of FIGS. 15 and 16 respectively, for example. One or more documents in the website are indexed using the indexing information received from one or more of the currently verified owners of the website (1704). At least a portion of the owners-information is provided to a respective one of the currently verified owners (1706). In response to a request from the respective one of the currently verified owners, the currently verified owners are re-verified (1708). Such a re-verification includes revoking the ownership rights of previously verified owners whose ownership rights have expired or terminated. In particular, if the verification tag or verification file of a "currently" verified owner (i.e., a previously verified owner) is not found by the server performing the verification process 1700, then the ownership rights of that owner are revoked, for example by marking the verification status 1514 (or 1614) of the user as "unverified". Typically, the verification tag or verification file of the previously verified owner would not be found because the verification tag or verification file has been deleted from the website.

FIG. 18 is a flowchart illustrating further details of the process for verifying ownership rights such as details associated with providing at least a portion of the owners-information to a respective one of the currently verified owners. The respective one of the currently verified owners is provided with the verification file and/or meta tag that exists for her account (1802). In other words, the respective one of the currently verified owners is provided with the owner-specific verification file and/or meta tag that is specific to her. Further, the respective one of the currently verified owners is provided with a complete list of verification files and/or meta tags that exist for the specified website (1804). The list includes verification files and/or meta tags that correspond to other currently verified owners of the specified website. The respective one of the currently verified owners can modify the list of verification files and/or meta tags to remove those verification files and/or meta tags that correspond to owners whose ownership rights have expired or terminated. The respective one of the currently verified owners can then request re-verification of owners based on the modified list of verification files and/or meta tags. The owners are re-verified and site owners for whom there is no verification file and/or meta tag are invalidated (1806).

Some non-limiting examples of mechanisms for enabling a user to revoke ownership rights and request re-verification of ownership rights regarding a website include a user interface, meta-tags stored at the website, and information stored in a predefined file stored in the website (e.g., a robots.txt file). As another non-limiting example, a predetermined protocol may be used to request re-verification of ownership rights regarding a website.

For purposes of explanation, assume that multiple webmasters M1, M2 and M3 have been verified previously as legitimate owners of a website W. Further assume that the ownership rights of M1 and M3 are terminated, subsequently. Search engines that are scheduled to crawl website W may be effectively notified of such termination of ownership rights by

19

the re-verification process as described herein. As a non-limiting example, webmaster M2 is provided, through an appropriate user interface, with a list of verification tags and/or verification files that currently exist for website W. Also, webmaster M2 is informed of the verification tag and/or verification file that is specific to webmaster M2. Webmaster M2 can use the user interface to remove verification tags and/or verification files that are no longer valid. After removal of the verification tags and/or verification files that are no longer valid, webmaster M2 can then request, through the user interface for example, a re-verification of the owners of the website W based on the updated list of verification tags and/or verification files.

According to certain embodiments, the system can periodically check the list of verification tags and/or verification files for updates without waiting for a website owner to request re-verification.

FIGS. 19 and 20 illustrate a non-limiting example of a GUI for enabling a website owner to manage site ownership verification. Site owner verification GUI 1900 of FIG. 19 includes information 1902 describing the use of the GUI for performing re-verification of website owners for a specified website, identification 1904 of one or more verification files and/or meta tags that are specific to a respective website owner who is using GUI 1900, a list 1906 of verification files and/or meta tags that exist for the specified website, and a re-verification button 1908. The website owner (also referred to as the “re-verifier”) who is attempting to perform the re-verification operation can edit list 1906 to remove verification files and/or meta tags that are no longer valid. Identification 1904 of the one or more verification files and/or meta tags that are specific to the re-verifier helps avoid the inadvertent removal of the re-verifier’s own verification files and/or meta tags during the re-verification process. When the re-verifier is satisfied that the invalid verification files and/or meta tags have been removed from list 1906, the re-verifier can activate the re-verification process by selecting re-verification button 1908. According to certain embodiments, upon re-verification, the re-verifier is shown GUI 2000 of FIG. 20. GUI 2000 includes re-verification information 2002, identification 2004 of one or more verification files and/or meta tags that are specific to the re-verifier, a list 2006 of verification and/or meta tags that exist for the specified website, and a re-verification button 2008. Re-verification information 2002 informs the re-verifier that the specified website has been re-verified, and may include information stating that: 1) site owners for whom a corresponding verification file and/or meta tag was not found are no longer verified for the specified website, and 2) site owners for whom a corresponding verification file and/or meta tag was found remain verified for the specified website. Identification 2004 of one or more verification files and/or meta tags that are specific to the re-verifier reminds the re-verifier of his or her verification files and/or meta tags to help the re-verifier review list 2006. In reviewing list 2006, if the re-verifier is satisfied that only valid site owners remain verified, then no further action is required of the re-verifier and the re-verifier may exit GUI 2000. If, however, the re-verifier finds that some invalid site owners remain verified, the re-verifier may once again edit list 2006. After editing list 2006, the re-verifier can re-activate the re-verification process by selecting re-verification button 2008.

According to certain embodiments, a site owner for a specified website may control the rate at which crawlers or crawl robots crawl the specified website. In some embodiments, the control of the crawl rate may be restricted to verified owners of the root domain of the website in order to limit the ability of sub-domain and virtual hosting site owners

20

to affect the portion of the website’s bandwidth used by web crawlers. FIG. 21 illustrates a crawl rate database 2100 used in controlling crawl rate, according to certain embodiments. As a non-limiting example, per site information database 740 of FIG. 7 may include crawl rate database 2100. Crawl rate database 2100 includes site crawl records 2102. Site crawl record 2102 includes site URL 2104, crawl rate limit 2106, crawl data 2108, a list of record pointers 2110, past utilization value 2112, and optionally, additional elements. Site URL 2104 stores the URL of the specified website. Crawl rate limit 2106 stores the crawl rate limit that either was originally pre-set by the web crawler system or that was subsequently specified by one of the owners of the specified website. For example, with the help of a crawl rate control module such as crawl rate control module of FIG. 7, an owner of the website may specify a crawl rate limit at which crawl robots 708 may crawl the specified website. As described in greater detail with reference to FIGS. 23, 24 and 25, the owner may or may not be allowed to increase the crawl rate limit depending on whether crawl rate limit is a factor in the crawling of the specified website. According to certain embodiments, crawl data 2108 includes statistical data 2122 on the number of web pages or documents accessed by the crawl robots at the specified website during a crawl session, statistical data 2124 on the time expended by the crawl robots on each access, and statistical data 2126 on the number bytes downloaded during the crawl session. List of log record pointers 2110 point to log records 2120 that include information related to accessing the web pages or documents during a crawl session. According to certain embodiments, the statistical data such as statistical data 2122, 2124 and 2126 are based on information stored in log records 2120. There may be numerous log records corresponding to crawl data captured over a period of time. Past utilization value 2112 is the highest crawl rate that occurred for the specified website over a pre-set period of time. In one embodiment, the pre-set period of time corresponds to a pre-set number of the most recent crawls of the specified website (e.g., the last 5 to 10 most recent crawls of the specified website). According to certain embodiments, the past utilization value 2112 may be the maximum number of pages accessed during a crawl session. Statistical data such as statistical data 2122, 2124 and 2126 and past utilization information such as past utilization value 2112 may be presented to an owner of the specified website through an appropriate user interface to aid in crawl rate control, as described in greater detail herein with reference to FIGS. 24 and 25.

FIG. 22 is a flowchart that describes a crawl rate control process 2200, according to certain embodiments. The specified website is crawled in accordance with a current rate limit (2202). For example, the current crawl rate limit may be a limit that is initially set by the website server, such as website server 200, for the specified website and subsequently re-set by an owner of the specified website. As previously described, crawl rate limit may be stored at crawl rate limit 2106 in crawl rate database 2100 of FIG. 21. Crawl data associated with the crawling of the website is stored (2204). For example, crawl data is stored in crawl data 2108 in crawl rate database 2100. A crawl rate control mechanism is presented to a respective owner of the specified website including at least a portion of the crawl data to enable the respective owner to select a new crawl rate limit (2206).

FIG. 23 is a flowchart that illustrates further details of crawl rate control as shown in crawl rate control process 2300. According to certain embodiments, the crawl rate control process includes determining if the crawl rate limit is a factor in crawling the specified website (2302). For example, crawl rate limit is a factor if the maximum number of documents

## 21

(past utilization value) accessed simultaneously by the crawl robots during a crawl session would be increased but for the crawl rate limit. In another example, the crawl rate limit is a factor if the difference between the current crawl rate limit and the past utilization value **2112** (or a corresponding rate value) is less than a predefined percentage (e.g., ten percent) of the current crawl rate limit. If crawl rate limit is a factor (**2302**—Yes), the website owner is allowed to select a faster crawl rate (**2304**), and as a result, the crawl rate limit for the specified website is increased (**2308**). If crawl rate limit is not a factor (**2302**—No), the website owner is informed that crawl rate limit is not a factor. According to certain embodiments, when crawl rate limit is not a factor, the website owner is not given an option to increase the crawl rate. According to certain other embodiments, when crawl rate limit is not a factor, the crawl rate limit is not increased despite the selection of a faster crawl rate by the website owner. In such a case, the website owner is informed that no change is made to the crawl rate limit. According to certain embodiments, the website owner can always choose a slower crawl rate. According to some embodiments, the change in crawl rate (either an increment or decrement in the crawl rate) as specified by the website owner may be applied immediately, including in the middle of a crawl session. In other embodiments, the change in crawl rate is applied during the next crawl session.

FIG. **24** is a schematic illustration of a crawl rate control GUI **2400** according to certain embodiments. GUI **2400** illustrates the case when crawl rate limit is not a factor in crawling the specified website. Crawl rate control GUI **2400** includes help information **2402** that describes the function of crawl rate in the context of collecting of documents from the specified website for purposes of indexing, for example. GUI **2400** also includes information **2404** to inform the website owner that crawl rate limit is not a factor in crawling the specified website. GUI **2400** further includes selection buttons such as a “faster” button **2406** for requesting a faster crawl rate if the website owner is allowed to select a faster crawl rate, a “normal” button **2408** which is a recommended crawl rate, a “slower” button **2410** for requesting a slower crawl rate, and a “save” button **2412** to save the website owner’s selection of crawl rate. GUI **2400** includes the presentation of a table **2414** and a graph **2416** of statistical data. For example, statistical data such as statistical data **2122**, **2124** and **2126** based on information stored in log records **2120** of FIG. **21** may be used to present data in table **2414** and graph **2416**. As a non-limiting example, graph **2416** shows the number of pages or documents accessed during a crawl session over a period of several months for the specified website. In some embodiments, projected statistical data is presented to the website owner based on the new crawl rate selected by the website owner. For example, projected statistical data includes expected bandwidth usage and/or expected number of pages or documents that may be crawled using the new crawl rate. In some embodiments, such projections may be estimated by revising recent statistical data referenced in FIG. **24** in a manner proportional to the specified change in crawl rate for the specified website.

Similarly, FIG. **25** is a schematic illustration of a crawl rate control GUI **2500**, according to certain embodiments. GUI **2500** illustrates the case when crawl rate limit is a factor in crawling the specified website. Crawl rate control GUI **2500** includes help information **2502** that describes the function of crawl rate in the context of collecting of documents from the specified website for purposes of indexing, for example. GUI **2500** also includes information **2504** to inform the website owner that the website owner may choose a faster crawl rate, if so desired. GUI **2400** further includes selection buttons

## 22

such as a “faster” button **2506** for requesting a faster crawl rate, a “normal” button **2508** which is a recommended crawl rate, a “slower” button **2510** for requesting a slower crawl rate, and a “save” button **2512** to save the website owner’s selection of crawl rate. GUI **2500** includes the presentation of a table **2514** and a graph **2516** of statistical data collected during a crawl session over a set period of time, for example.

The foregoing description, for purpose of explanation, has been described with reference to specific embodiments. However, the illustrative discussions above are not intended to be exhaustive or to limit the invention to the precise forms disclosed. Many modifications and variations are possible in view of the above teachings. The embodiments were chosen and described in order to best explain the principles of the invention and its practical applications, to thereby enable others skilled in the art to best utilize the invention and various embodiments with various modifications as are suited to the particular use contemplated.

We claim:

**1.** A method of indexing documents in websites, the method comprising:

on a server system having one or more processors and memory storing programs to be executed by the one or more processors:

storing, for each website of a multiplicity of websites, a corresponding current crawl rate limit;

comparing a maximum crawl rate for a respective website over a defined period of time with the current crawl rate limit for crawling the respective website to determine if the current crawl rate limit is a limiting factor in crawling the respective website;

performing a website crawling management function in accordance with the determination of whether the current crawl rate limit is the limiting factor in crawling the respective website; and

providing a crawl rate control mechanism to a respective owner of the respective website, wherein the crawl rate control mechanism enables selection of a new crawl rate limit corresponding to the respective website by the respective owner.

**2.** The method of claim **1**, further including crawling the respective website at a rate no greater than the current crawl rate limit for crawling the respective website.

**3.** The method of claim **1**, further comprising:

when the current crawl rate limit is not the limiting factor in crawling the respective website, informing the respective owner that a request for a faster crawl rate may not change a current crawl rate for crawling the respective website.

**4.** The method of claim **1**, wherein the current crawl rate limit is the limiting factor only when a difference between the current crawl rate limit and the maximum crawl rate for the respective website over the defined period of time is less than a predefined quantity.

**5.** The method of claim **1**, further comprising:

in response to a request to decrease the current crawl rate limit for crawling the respective website, decreasing the current crawl rate limit for crawling the respective website.

**6.** The method of claim **1**, further comprising determining a number of documents of the respective website that are accessed during one or more crawl sessions and storing corresponding crawl data.

**7.** The method of claim **6**, wherein storing the corresponding crawl data further comprises determining an average quantity of time expended to access the documents from the respective website during the one or more crawl sessions.

23

8. The method of claim 1, further comprising determining a number of bytes downloaded from the respective website during one or more crawl sessions and storing corresponding crawl data.

9. The method of claim 1, including providing, for display, resource usage statistics corresponding to resources of the respective website used during a plurality of prior crawl visits of the respective website.

10. The method of claim 1, wherein providing the crawl rate mechanism includes providing, for concurrent display: the current crawl rate limit associated with the crawling of the respective website; crawl data, including statistical information associated with crawling the respective website; and an interface for enabling the respective owner to select the new crawl rate limit.

11. The method of claim 10, wherein the providing for concurrent display includes providing for display recommendations for selecting the new crawl rate limit, wherein the recommendations are based, at least in part, on whether the current crawl rate limit is the limiting factor in crawling the respective website.

12. A computer system comprising: memory; one or more processors; and at least one program stored in the memory and executed by the one or more processors, the at least one program including instructions for: storing, for each website of a multiplicity of websites, a corresponding current crawl rate limit; comparing a maximum crawl rate for a respective website over a defined period of time with the current crawl rate limit for crawling the respective website to determine if the current crawl rate limit is a limiting factor in crawling the respective website; performing a website crawling management function in accordance with the determination of whether the current crawl rate limit is the limiting factor in crawling the respective website; and providing a crawl rate control mechanism to a respective owner of the respective website, wherein the crawl rate control mechanism enables selection of a new crawl rate limit corresponding to the respective website by the respective owner.

13. The computer system of claim 12, the at least one program further including instructions for crawling the respective website at a rate no greater than the current crawl rate limit for crawling the respective website.

14. The computer system of claim 12, the at least one program further comprising: instructions for informing the respective owner that a request for a faster crawl rate may not change a current crawl rate for crawling the respective website when the current crawl rate limit is not the limiting factor in crawling the respective website.

15. The computer system of claim 12, wherein the current crawl rate limit is the limiting factor only when a difference between the current crawl rate limit and the maximum crawl rate for the respective website over the defined period of time is less than a predefined quantity.

16. The computer system of claim 12, the instructions further comprising: instructions for decreasing the current crawl rate limit for crawling the respective website in response to a request to decrease the current crawl rate limit for crawling the respective website.

24

17. The computer system of claim 12, wherein the at least one program further comprises instructions for determining a number of documents of the respective website that are accessed during one or more crawl sessions and for storing corresponding crawl data.

18. The computer system of claim 17, wherein the instructions for storing the corresponding crawl data further comprise instructions for determining an average quantity of time expended to access the documents from the respective website during the one or more crawl sessions.

19. The computer system of claim 12, wherein the at least one program further comprises instructions for determining a number of bytes downloaded from the respective website during one or more crawl sessions and for storing corresponding crawl data.

20. The computer system of claim 12, wherein the at least one program further comprises instructions for providing, for display, resource usage statistics corresponding to resources of the respective website used during a plurality of prior crawl visits of the respective website.

21. The computer system of claim 12, wherein the instructions for providing the crawl rate control mechanism further include instructions providing, for concurrent display: the current crawl rate limit associated with the crawling of the respective website; crawl data, including statistical information associated with crawling the respective website; and an interface for enabling the respective owner to select the new crawl rate limit.

22. The computer system of claim 21, wherein the instructions for providing for concurrent display further include instructions for providing an information display having recommendations for selecting the new crawl rate limit, wherein the recommendations are based, at least in part, on whether the current crawl rate is the limiting factor in crawling the website.

23. A non-transitory computer readable storage medium storing at least one program for execution by a computer system, the at least one program comprising instructions for: storing, for each website of a multiplicity of websites, a corresponding current crawl rate limit; comparing a maximum crawl rate for the respective website over a defined period of time with the current crawl rate limit for crawling the respective website to determine if the current crawl rate limit is a limiting factor in crawling the respective website; and performing a website crawling management function in accordance with the determination of whether the current crawl rate limit is the limiting factor in crawling the respective website; and providing a crawl rate control mechanism to a respective owner of the respective website, wherein the crawl rate control mechanism enables selection of a new crawl rate limit corresponding to the respective website by the respective owner.

24. The non-transitory computer program product of claim 23, the at least one program further including instructions for crawling the respective website at a rate no greater than the current crawl rate limit for crawling the respective website.

25. The non-transitory computer program product of claim 23, the at least one program further comprising: instructions for informing the respective owner that a request for a faster crawl rate may not change a current crawl rate for crawling the respective website when the current crawl rate limit is not the limiting factor in crawling the respective website.

25

26. The non-transitory computer program product of claim 23, wherein the current crawl rate limit is the limiting factor only when a difference between the current crawl rate limit and the maximum crawl rate for the respective website over the defined period of time is less than a predefined quantity. 5

27. The non-transitory computer program product of claim 23, wherein the at least one program further comprises: instructions for decreasing the current crawl rate limit in response to a request to decrease the current crawl rate limit for crawling the respective website for crawling the 10  
respective website.

28. The non-transitory computer program product of claim 23, wherein the at least one program further comprises instructions for determining a number of documents of the respective website that are accessed during one or more crawl sessions and for storing corresponding crawl data. 15

29. The non-transitory computer program product of claim 28, wherein the instructions for storing the corresponding crawl data further comprise instructions for determining an average quantity of time expended to access the documents from the respective website during the one or more crawl sessions. 20

30. The non-transitory computer program product of claim 23, wherein the instructions for storing crawl data further comprise instructions for determining a number of bytes downloaded from the respective website during one or more crawl sessions and for storing corresponding crawl data. 25

31. The non-transitory computer program product of claim 23, wherein the at least one program further comprises instructions for providing, for display, resource usage statistics corresponding to resources of the respective website used during a plurality of prior crawl visits of the website. 30

32. The non-transitory computer program product of claim 23, wherein the instructions for providing the crawl rate control mechanism further include instructions providing, for concurrent display: 35

the current crawl rate associated with the crawling of the respective website;

26

crawl data, including statistical information associated with crawling the respective website; and  
an interface for enabling the respective owner to select the new crawl rate limit.

33. The non-transitory computer program product of claim 32, wherein the instructions for providing for concurrent display further include instructions for providing for display recommendations for selecting the new crawl rate limit, wherein the recommendations are based, at least in part, on whether the current crawl rate limit is the limiting factor in crawling the respective website.

34. The method of claim 1, including providing, for display, a status log indicating respective statuses of documents of the respective website that are accessed during one or more crawl sessions.

35. The method of claim 1, wherein storing the corresponding crawl rate limit further comprises storing a list of respective time intervals and a crawl rate limit for each time interval.

36. The computer system of claim 12, the at least one program further including instructions for providing, for display, a status log indicating respective statuses of documents of the respective website that are accessed during one or more crawl sessions.

37. The computer system of claim 12, the instructions for storing the corresponding crawl rate limit further comprising storing a list of respective time intervals and a crawl rate limit for each time interval.

38. The non-transitory computer readable storage medium of claim 23, the at least one program further including instructions for providing, for display, a status log indicating respective statuses of documents of the respective website that are accessed during one or more crawl sessions.

39. The non-transitory computer readable storage medium of claim 23, the instructions for storing the corresponding crawl rate limit further comprising storing a list of respective time intervals and a crawl rate limit for each time interval.

\* \* \* \* \*